

Ray: A Distributed Framework for Emerging AI Applications & Lineage Stash: Fault Tolerance Off the Critical Path

Shucheng Zhong, Han You, Joshua Segal

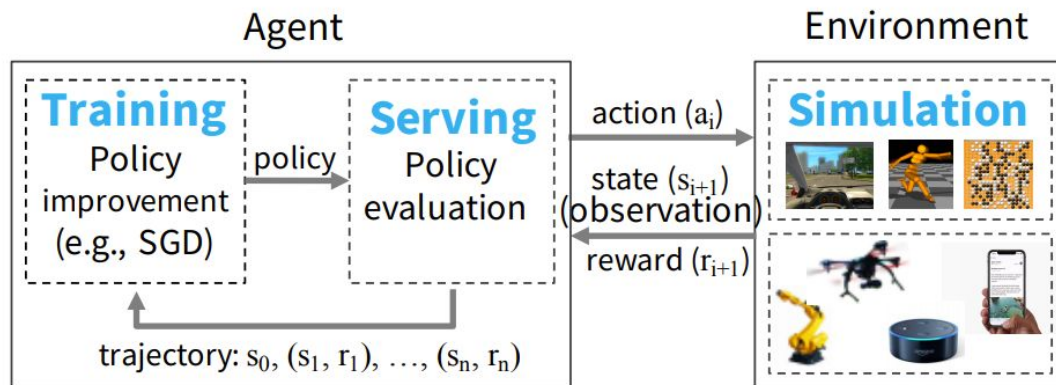
Overview

- Ray
 - Background & Motivation
 - API
 - Architecture
- Lineage Stash
 - Background & Motivation
 - Architecture
- Evaluations of Ray and Lineage Stash

Systems for reinforcement learning

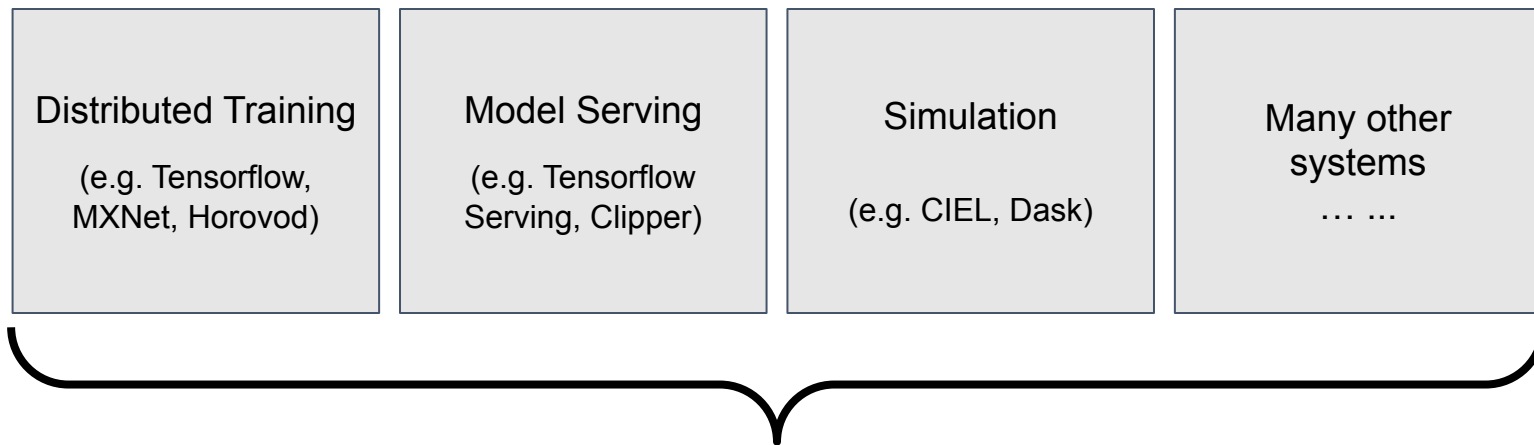
Three parts of RL:

- Simulation
- Training
- Serving



* Figure from *Ray: A Distributed Framework for Emerging AI Applications*

Machine learning ecosystem nowadays

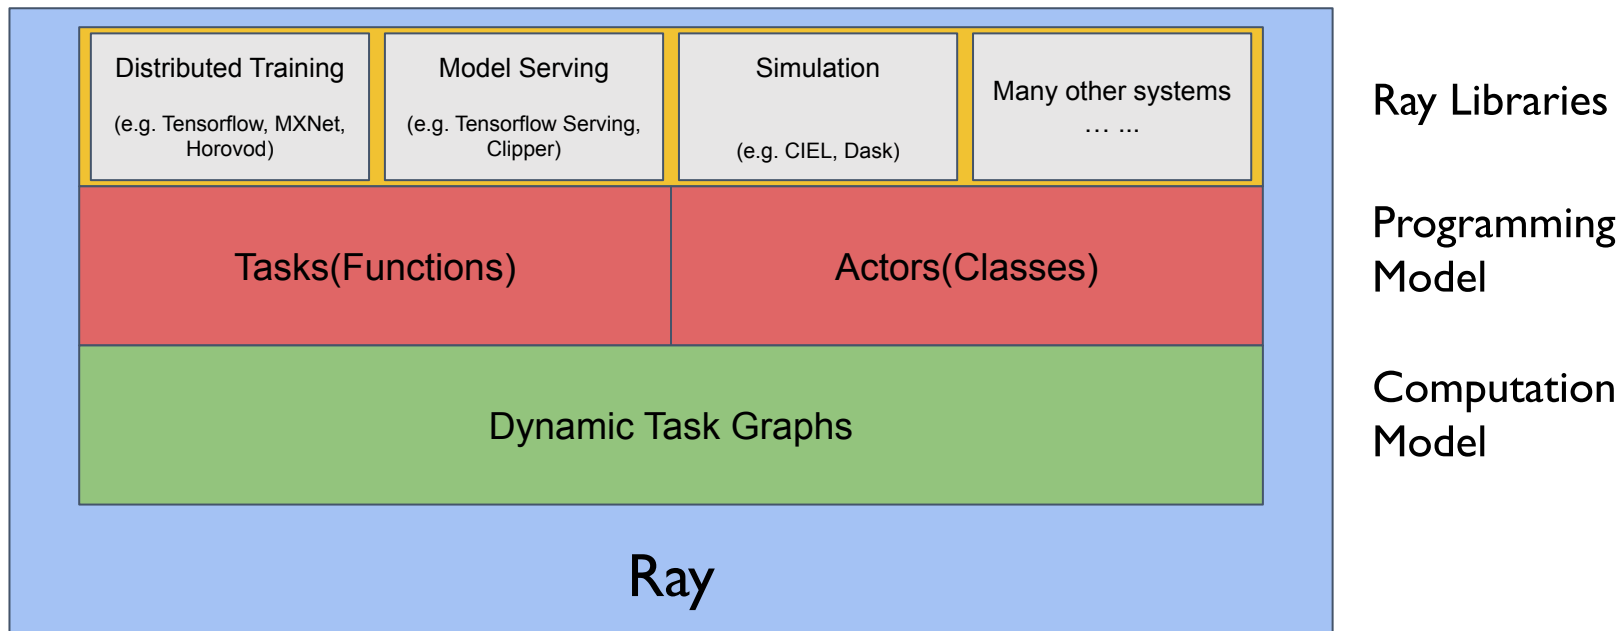


What is Ray?

A general-purpose cluster-computing framework that supports

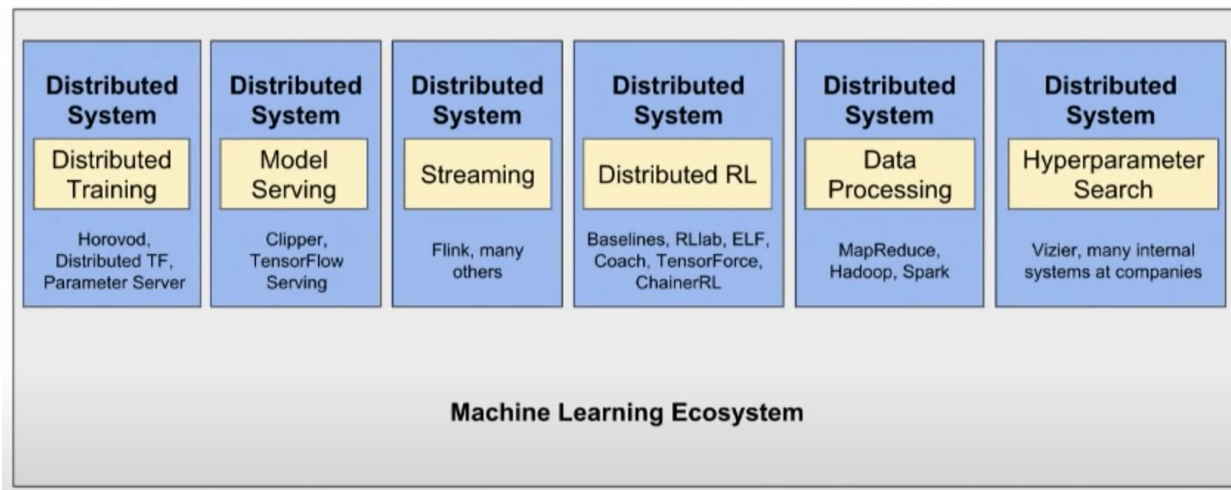
- fine-grained computations (millions of tasks within millisecond-level latency)
- Heterogeneity both in time and in resource usage
- dynamic execution

Overview of Ray



Motivation


The Old Way: Glueing a bunch of systems together -- or build from scratch




Ray's Computation Model

Dependencies represented as edges in task graph.

```
@ray.remote  
def func():  
    foo_result = foo.remote()  
    bar_result = bar.remote(foo_result)
```




Control dependency



Dataflow dependency

Ray's Computation Model

Tasks and Actors are automatically triggered when inputs become available.



Ray API: Tasks

```
@ray.remote
def matrixMut(A, B):
    return A * B

@ray.remote
def matrixFunction(A, B):
    return np.eig(A) * np.eig(B)

@ray.remote
def matrixSum(A, B):
    return A + B
```

$$A \times B + f(B, C)$$

```
AxB = matrixMut.remote(A, B)
f_BC = matrixFunction.remote(B, C)
result = matrixSum.remote(AxB, f_BC)
print ray.get(result)
```

Ray API: Actor

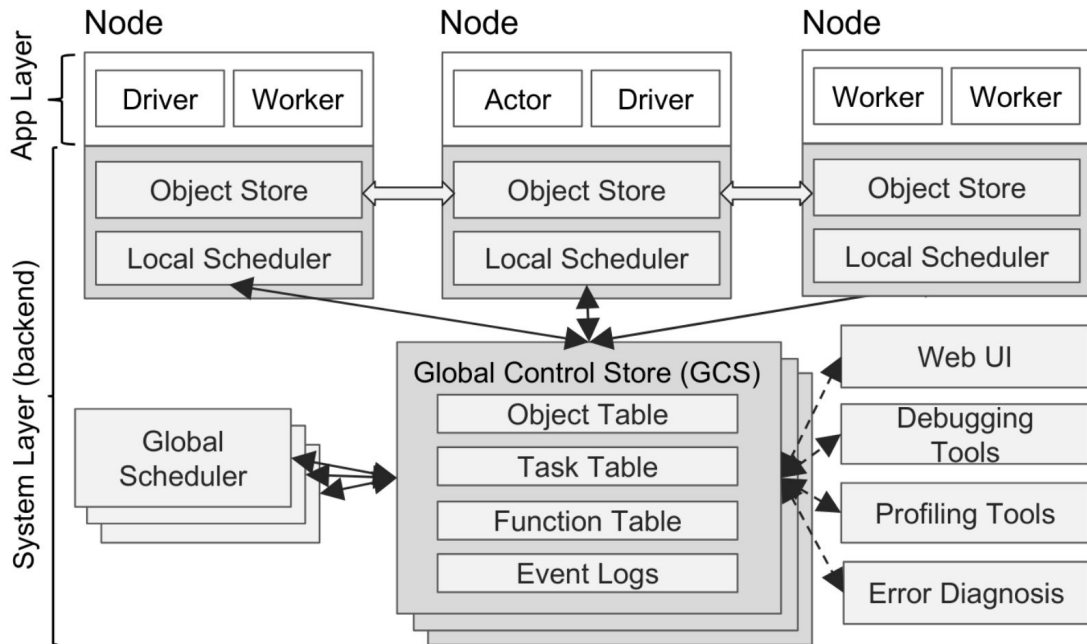
```
@ray.remote
class Shuffler:
    def __init__(self, A):
        self.A = A

    def shuffle(self, seed):
        self.A = shuffle(self.A, seed)
        return self.A
```


```
shuffler = Shuffler.remote(result)
iter1 = shuffler.shuffle.remote(seed)
iter2 = shuffler.shuffle.remote(seed)
iter3 = shuffler.shuffle.remote(seed)
```

$A' = \text{shuffle}(A, s)$


Architecture




$$A \times B + f(B, C)$$




Node A




Node B




$$A \times B + f(B, C)$$



Node A




Node B




→

```
AxB = matrixMut.remote(A, B)
f_BC = matrixFunction.remote(B, C)
result = matrixSum.remote(AxB, f_BC)
print ray.get(result)
```


$$A \times B + f(B, C)$$



Node A



Node B




Local scheduler of A: Let's do this ourselves

➔

```
AxB = matrixMut.remote(A, B)
f_BC = matrixFunction.remote(B, C)
result = matrixSum.remote(AxB, f_BC)
print ray.get(result)
```


$$A \times B + f(B, C)$$



Local scheduler of A: Let's do this ourselves

```
➔ AxB = matrixMut.remote(A, B)
  f_BC = matrixFunction.remote(B, C)
  result = matrixSum.remote(AxB, f_BC)
  print ray.get(result)
```



$$A \times B + f(B, C)$$



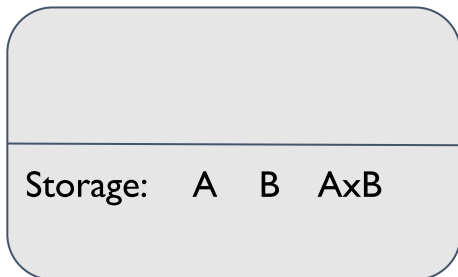
Local scheduler of A: Let's do this ourselves

```
➡ AxB = matrixMut.remote(A, B)
  f_BC = matrixFunction.remote(B, C)
  result = matrixSum.remote(AxB, f_BC)
  print ray.get(result)
```


$$A \times B + f(B, C)$$



Node A



Node B




Local scheduler of A: Let's do this ourselves


➔

```
AxB = matrixMut.remote(A, B)
f_BC = matrixFunction.remote(B, C)
result = matrixSum.remote(AxB, f_BC)
print ray.get(result)
```


$$A \times B + f(B, C)$$



Node A




Node B

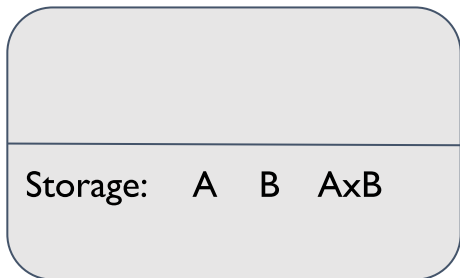


```
AxB = matrixMut.remote(A, B)  
f_BC = matrixFunction.remote(B, C)  
result = matrixSum.remote(AxB, f_BC)  
print ray.get(result)
```


$$A \times B + f(B, C)$$



Node A




Node B



Local scheduler of A: I don't want to do this

```
AxB = matrixMut.remote(A, B)  
f_BC = matrixFunction.remote(B, C)  
result = matrixSum.remote(AxB, f_BC)  
print ray.get(result)
```


$$A \times B + f(B, C)$$




Local scheduler of A: I don't want to do this

```
→ AxB = matrixMut.remote(A, B)
f_BC = matrixFunction.remote(B, C)
result = matrixSum.remote(AxB, f_BC)
print ray.get(result)
```


$$A \times B + f(B, C)$$



Node A




Node B




Local scheduler of A: I don't want to do this

```
AxB = matrixMut.remote(A, B)
f_BC = matrixFunction.remote(B, C)
result = matrixSum.remote(AxB, f_BC)
print ray.get(result)
```


$$A \times B + f(B, C)$$



Node A




Node B




```
AxB = matrixMut.remote(A, B)  
f_BC = matrixFunction.remote(B, C)  
result = matrixSum.remote(AxB, f_BC)  
print ray.get(result)
```

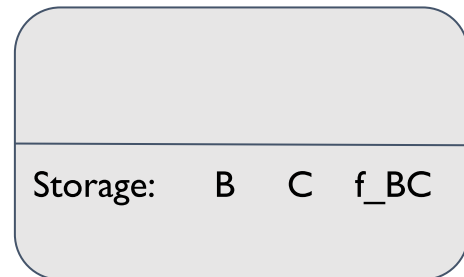
$$A \times B + f(B, C)$$



Node A




Node B




```
AxB = matrixMut.remote(A, B)
f_BC = matrixFunction.remote(B, C)
→ result = matrixSum.remote(AxB, f_BC)
print ray.get(result)
```



$$A \times B + f(B, C)$$



Node A



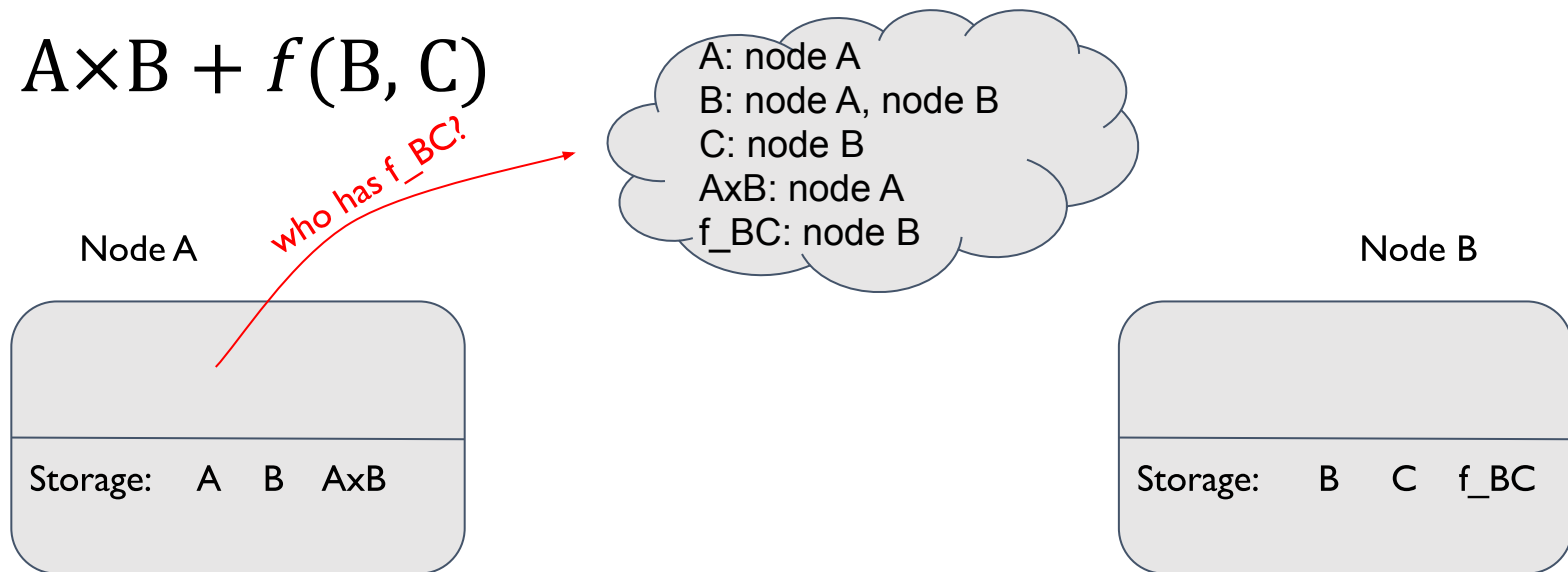
Node B



Local scheduler of A: Let's do this ourselves

```
AxB = matrixMut.remote(A, B)
f_BC = matrixFunction.remote(B, C)
→ result = matrixSum.remote(AxB, f_BC)
print ray.get(result)
```

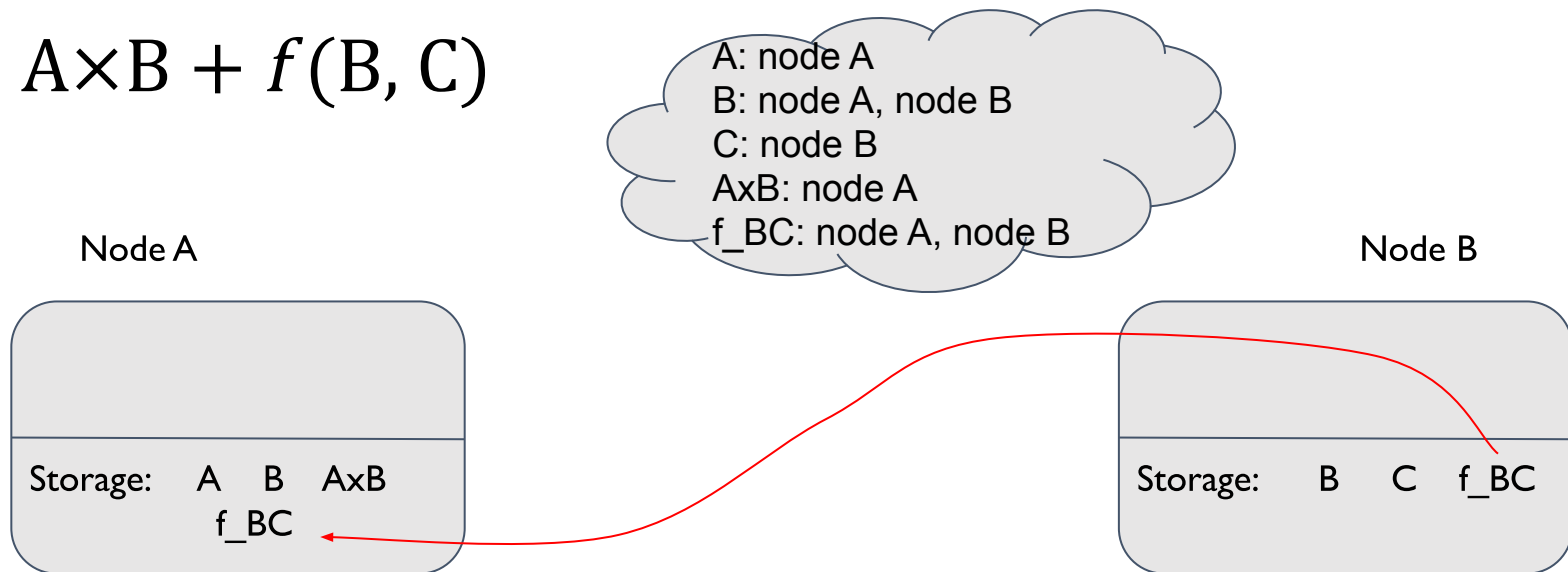
$$A \times B + f(B, C)$$



Local scheduler of A: Let's do this ourselves

```
AxB = matrixMut.remote(A, B)
f_BC = matrixFunction.remote(B, C)
result = matrixSum.remote(AxB, f_BC)
print ray.get(result)
```


$$A \times B + f(B, C)$$



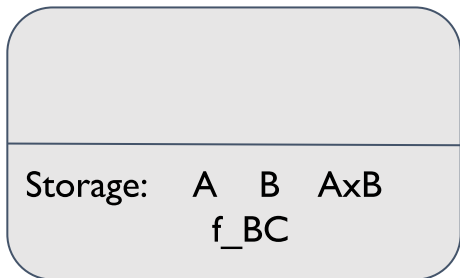
Local scheduler of A: Let's do this ourselves

```
AxB = matrixMut.remote(A, B)
f_BC = matrixFunction.remote(B, C)
result = matrixSum.remote(AxB, f_BC)
print ray.get(result)
```

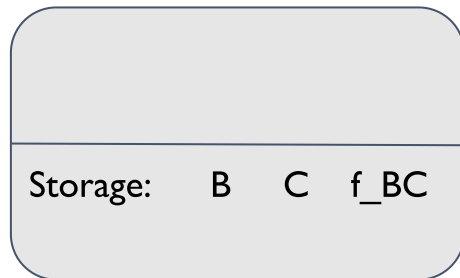
$$A \times B + f(B, C)$$



Node A




Node B



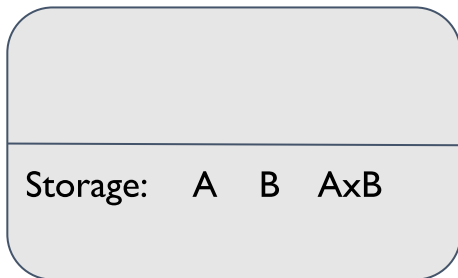
Local scheduler of A: Let's do this ourselves

```
AxB = matrixMut.remote(A, B)
f_BC = matrixFunction.remote(B, C)
result = matrixSum.remote(AxB, f_BC)
print ray.get(result)
```


$$A \times B + f(B, C)$$



Node A



Node B




Local scheduler of A: Let's do this ourselves

```
AxB = matrixMut.remote(A, B)
f_BC = matrixFunction.remote(B, C)
→ result = matrixSum.remote(AxB, f_BC)
print ray.get(result)
```

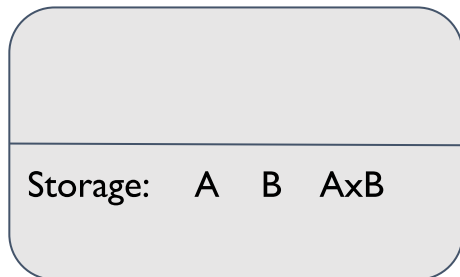


Break!


$$A \times B + f(B, C)$$



Node A



Node B



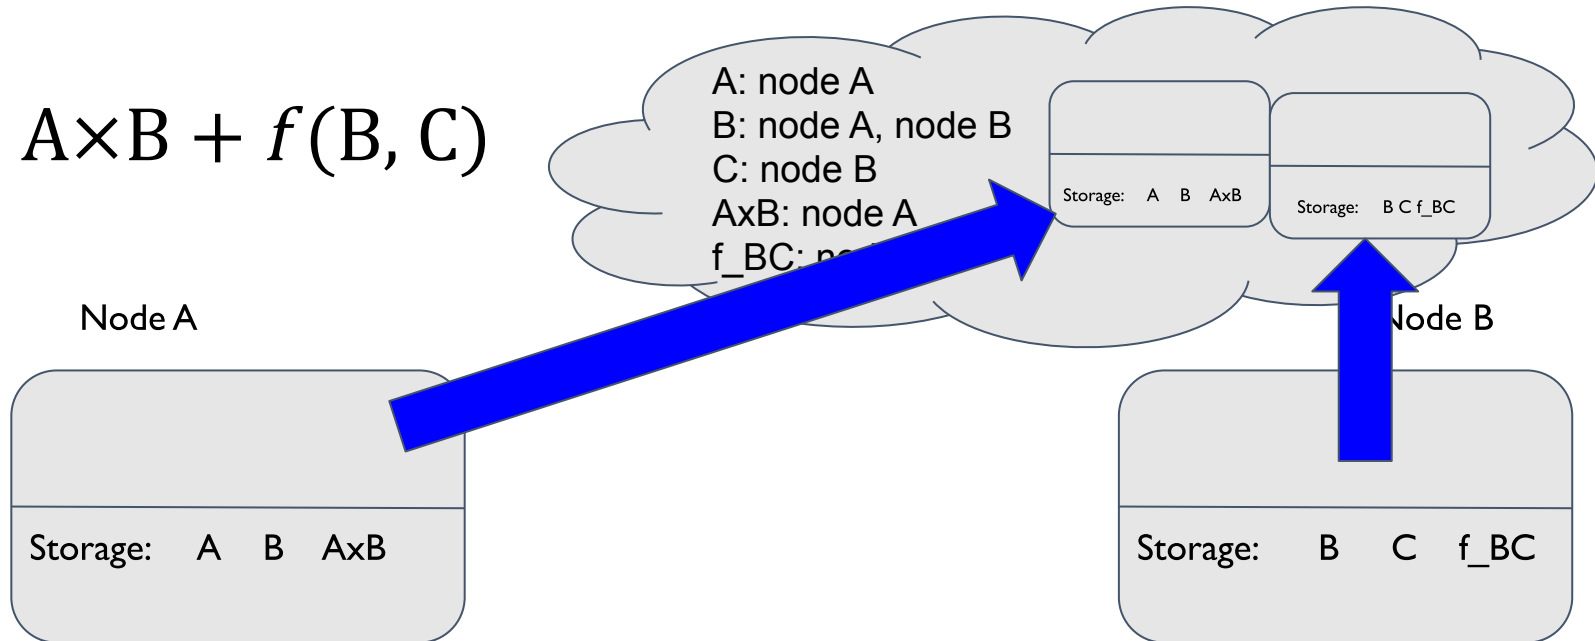
Local scheduler of A: Let's do this ourselves

```
AxB = matrixMut.remote(A, B)
f_BC = matrixFunction.remote(B, C)
→ result = matrixSum.remote(AxB, f_BC)
print ray.get(result)
```


Attempt 1: Global Checkpoint

- save the entire state of the workers periodically

$$A \times B + f(B, C)$$


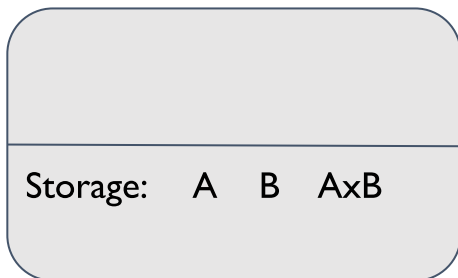
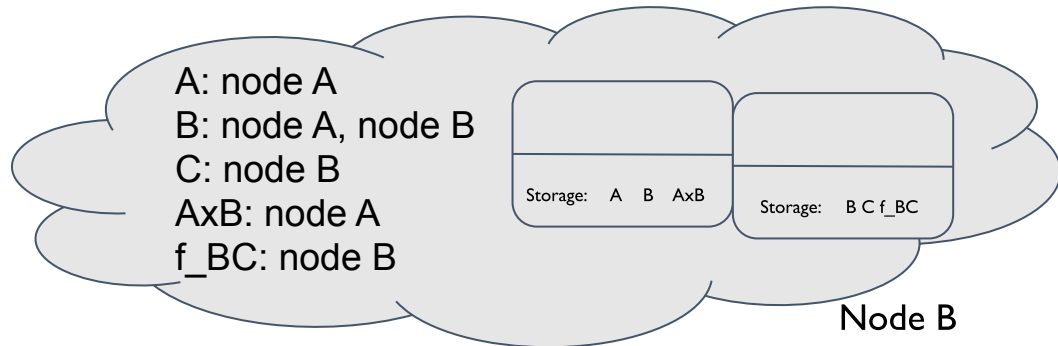


Local scheduler of A: Let's do this ourselves

```

AxB = matrixMut.remote(A, B)
f_BC = matrixFunction.remote(B, C)
result = matrixSum.remote(AxB, f_BC)
print ray.get(result)
  
```


$$A \times B + f(B, C)$$



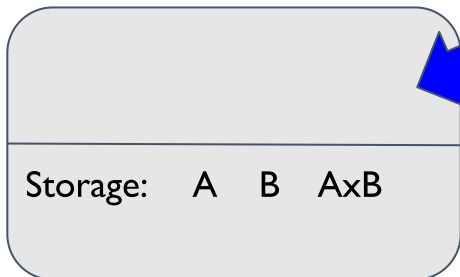
Local scheduler of A: Let's do this ourselves

```
AxB = matrixMut.remote(A, B)
f_BC = matrixFunction.remote(B, C)
result = matrixSum.remote(AxB, f_BC)
print ray.get(result)
```


$$A \times B + f(B, C)$$



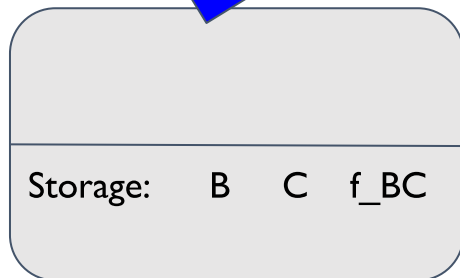
Node A



Node B



Node C



Local scheduler of A: Let's do this ourselves

```

AxB = matrixMut.remote(A, B)
f_BC = matrixFunction.remote(B, C)
result = matrixSum.remote(AxB, f_BC)
print ray.get(result)
  
```

Global Checkpoint

- high recovery time

Global Checkpoint

- high recovery time
- low overhead
 - State is sent in the background asynchronously

Global Checkpoint

- high recovery time
- low overhead
 - State is sent in the background asynchronously
- good for fine grained tasks (millisecond)
 - small batches
 - streams of data


Attempt 2: Lineage Logging

- Store the lineage


Attempt 2: Lineage Logging

- Save the logs on what's dependent on what

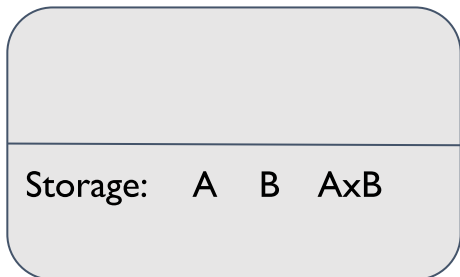
```
AxB = matrixMut.remote(A, B)
f_BC = matrixFunction.remote(B, C)
result = matrixSum.remote(AxB, f_BC)
print ray.get(result)
```



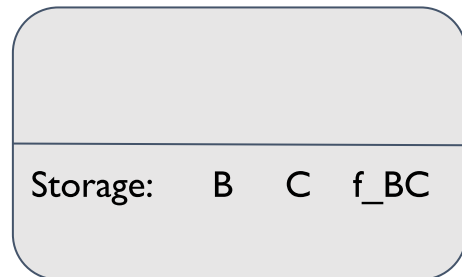
$$A \times B + f(B, C)$$



Node A




Node B




Local scheduler of A: Let's do this ourselves

```
AxB = matrixMut.remote(A, B)
f_BC = matrixFunction.remote(B, C)
result = matrixSum.remote(AxB, f_BC)
print ray.get(result)
```


$$A \times B + f(B, C)$$



Node A



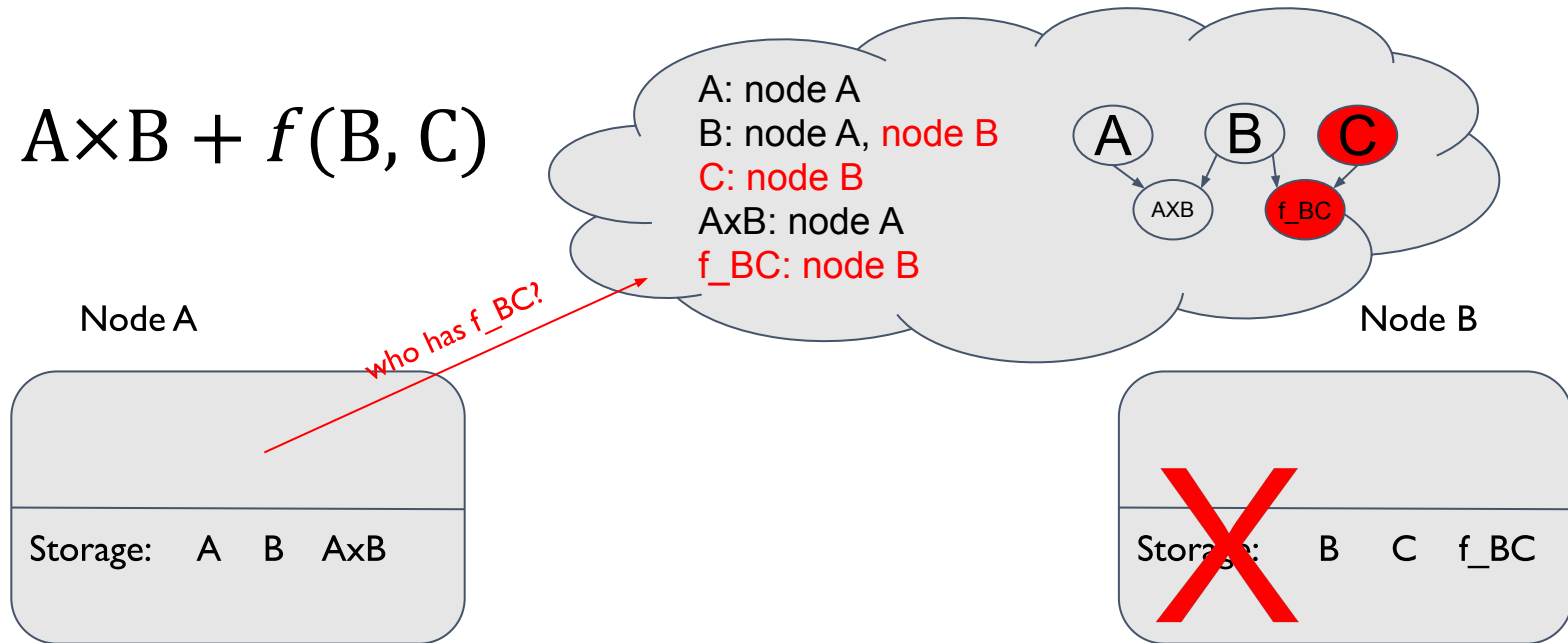
Node B



Local scheduler of A: Let's do this ourselves

```
AxB = matrixMut.remote(A, B)
f_BC = matrixFunction.remote(B, C)
result = matrixSum.remote(AxB, f_BC)
print ray.get(result)
```

$$A \times B + f(B, C)$$



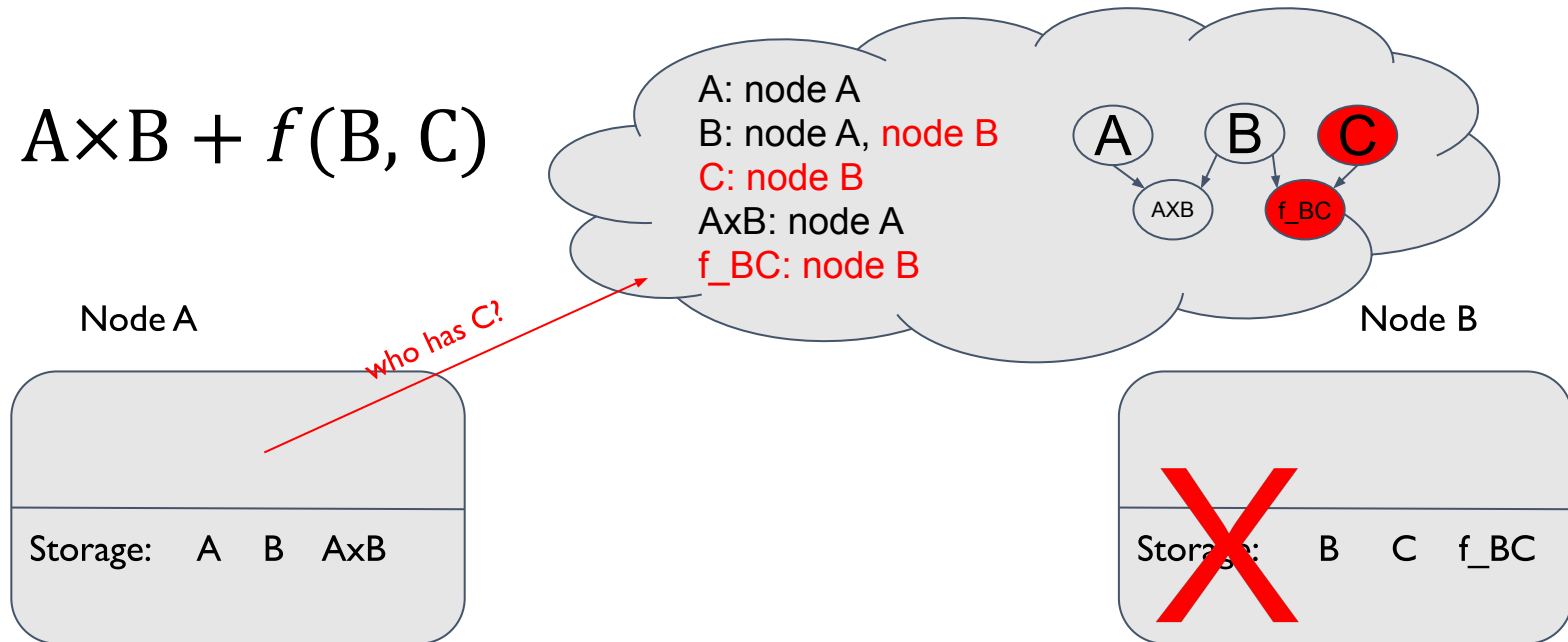
Local scheduler of A: Let's do this ourselves

```

AXB = matrixMut.remote(A, B)
f_BC = matrixFunction.remote(B, C)
result = matrixSum.remote(AXB, f_BC)
print ray.get(result)

```

$$A \times B + f(B, C)$$




Local scheduler of A: Let's do this ourselves

```

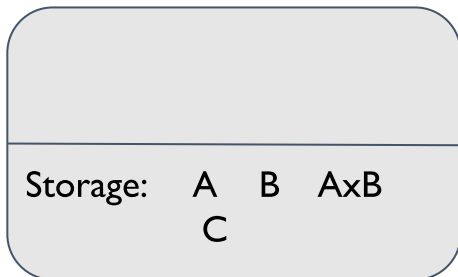
AXB = matrixMut.remote(A, B)
f_BC = matrixFunction.remote(B, C)
result = matrixSum.remote(AXB, f_BC)
print ray.get(result)

```


$$A \times B + f(B, C)$$



Node A




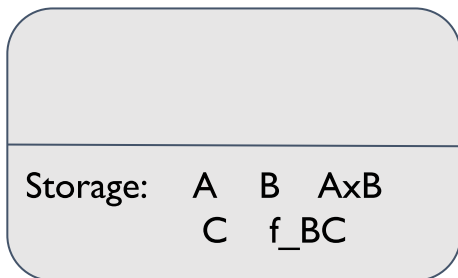

Node B



Local scheduler of A: Let's do this ourselves

```
AxB = matrixMut.remote(A, B)
f_BC = matrixFunction.remote(B, C)
result = matrixSum.remote(AxB, f_BC)
print ray.get(result)
```

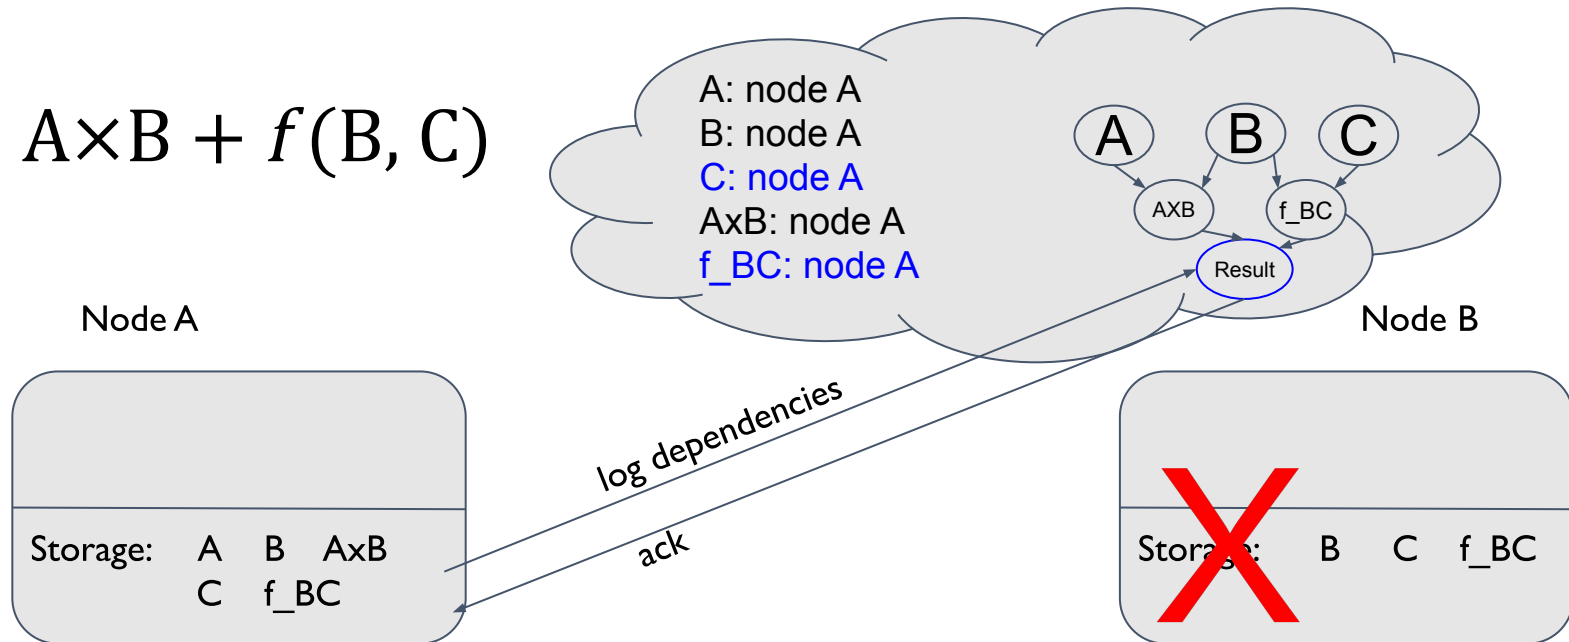
$$A \times B + f(B, C)$$



Local scheduler of A: Let's do this ourselves

```
AxB = matrixMut.remote(A, B)
f_BC = matrixFunction.remote(B, C)
result = matrixSum.remote(AxB, f_BC)
print ray.get(result)
```

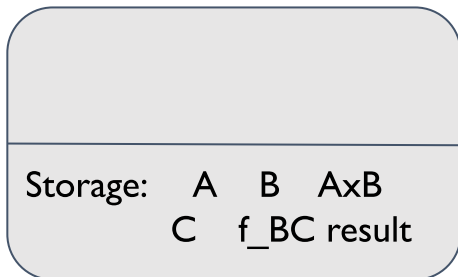
$$A \times B + f(B, C)$$




```
AxB = matrixMut.remote(A, B)
f_BC = matrixFunction.remote(B, C)
→ result = matrixSum.remote(AxB, f_BC)
print ray.get(result)
```


$$A \times B + f(B, C)$$


Node A



A: node A
 B: node A
 C: node A
 AxB: node A
 f_BC: node A
 Result: node A




Node B



```
AxB = matrixMut.remote(A, B)
f_BC = matrixFunction.remote(B, C)
result = matrixSum.remote(AxB, f_BC)
print ray.get(result)
```

$$A \times B + f(B, C)$$


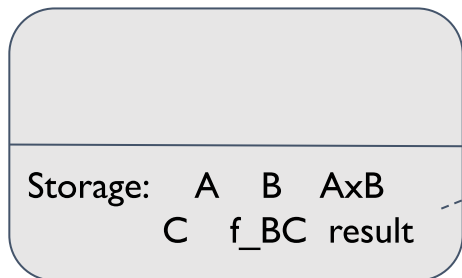



```

AxB = matrixMut.remote(A, B)
f_BC = matrixFunction.remote(B, C)
→ result = matrixSum.remote(AxB, f_BC)
print ray.get(result)

```

$$A \times B + f(B, C)$$




Local scheduler of A: I don't want to do this


```
AxB = matrixMut.remote(A, B)
f_BC = matrixFunction.remote(B, C)
result = matrixSum.remote(AxB, f_BC)
```

→ result2 = matrixSum(result, A)


$$A \times B + f(B, C)$$



Node A



Node B




Local scheduler of A: I don't want to do this


```
AxB = matrixMut.remote(A, B)
f_BC = matrixFunction.remote(B, C)
result = matrixSum.remote(AxB, f_BC)
```

→ result2 = matrixSum(result, A)


Node C



$$A \times B + f(B, C)$$




Node A



log dependencies

Node B




Local scheduler of A: I don't want to do this


```
AxB = matrixMut.remote(A, B)
f_BC = matrixFunction.remote(B, C)
result = matrixSum.remote(AxB, f_BC)
```

→ result2 = matrixSum(result, A)


Node C



$$A \times B + f(B, C)$$




Node A



X log dependencies

Node B




Local scheduler of A: I don't want to do this

```
AxB = matrixMut.remote(A, B)
f_BC = matrixFunction.remote(B, C)
result = matrixSum.remote(AxB, f_BC)
```

→ result2 = matrixSum(result, A)

Node C



Lineage Logging

Lineage Logging

- Fast recovery time
- Large overhead
 - need to commit lineage before each task
 - need to wait for ack before starting each task

Lineage Logging

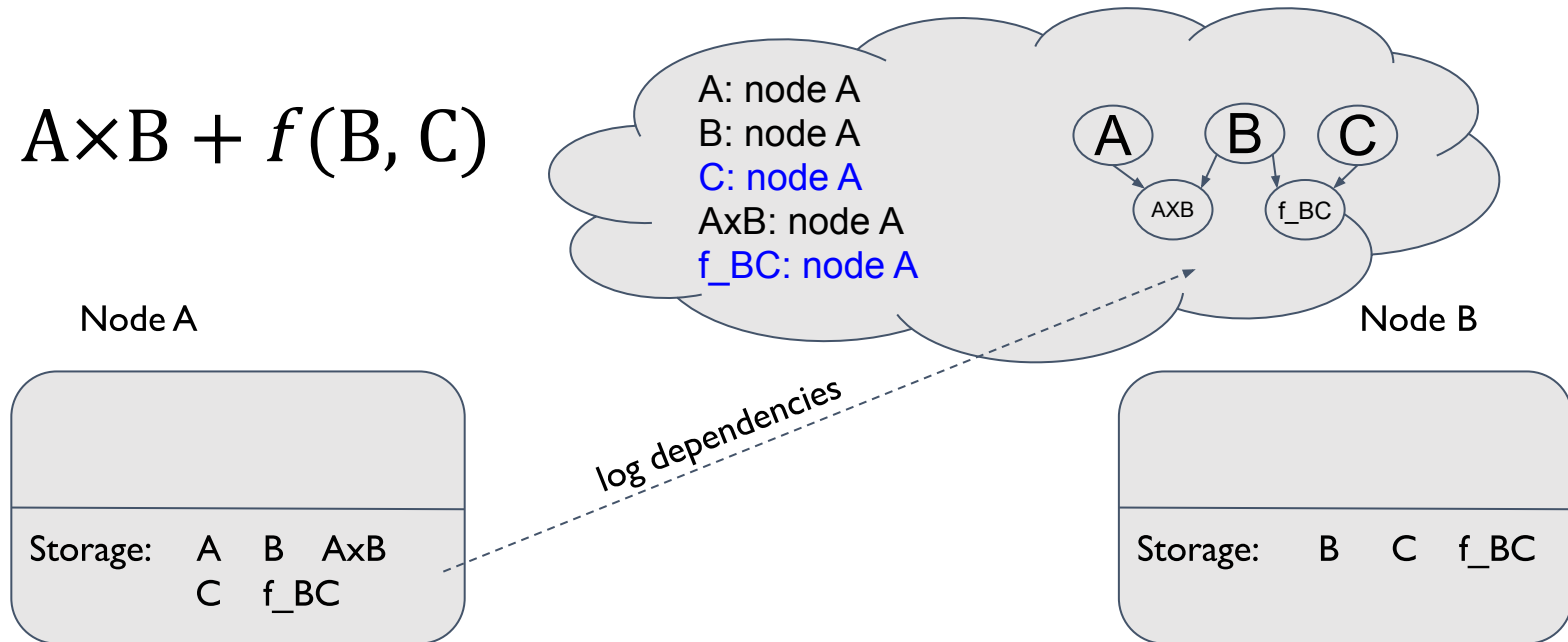
- Fast recovery time
- Large overhead
 - need to commit lineage before each task
 - need to wait for ack before starting each task
- Good for coarse grained tasks (seconds)
 - Big data

Lineage Stashing

Lineage Stashing

- low recovery time
- low overhead

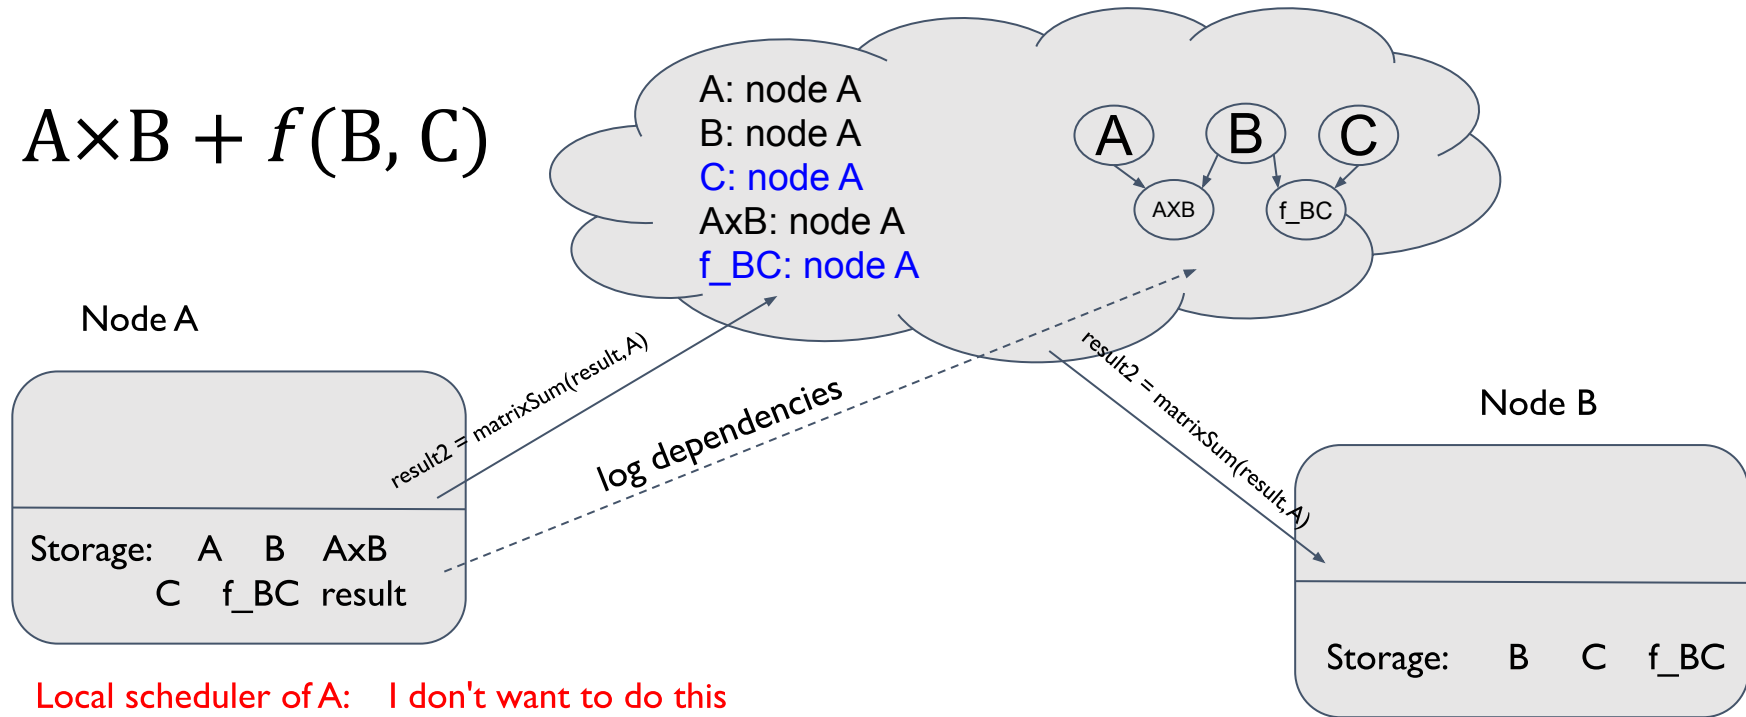
$$A \times B + f(B, C)$$



```

AxB = matrixMut.remote(A, B)
f_BC = matrixFunction.remote(B, C)
→ result = matrixSum.remote(AxB, f_BC)
    result2 = matrixSum(result, A)
  
```

$$A \times B + f(B, C)$$

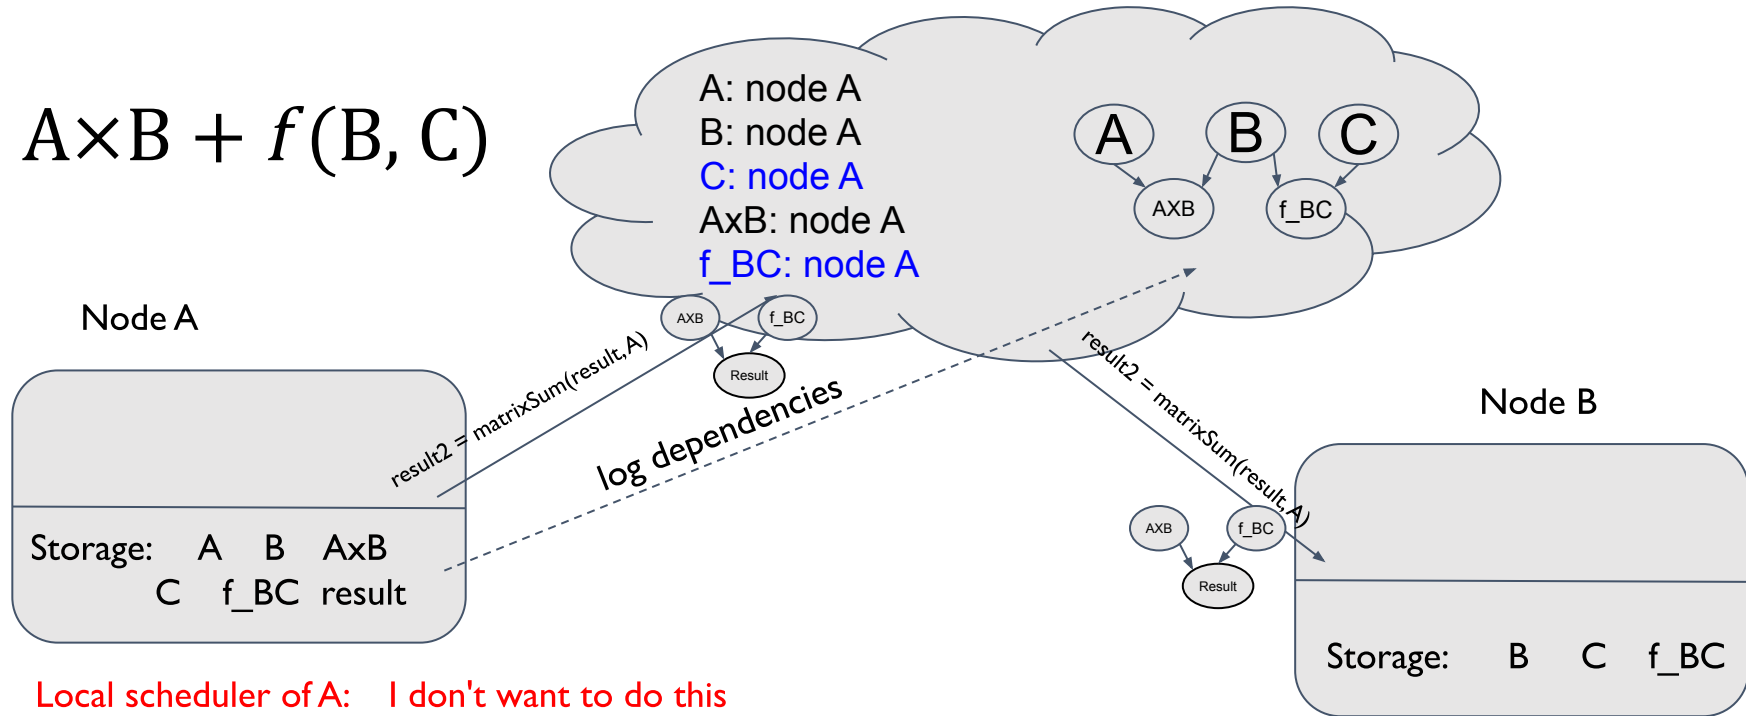


Local scheduler of A: I don't want to do this

```
AxB = matrixMut.remote(A, B)
f_BC = matrixFunction.remote(B, C)
result = matrixSum.remote(AxB, f_BC)
```

→ result2 = matrixSum(result, A)

$$A \times B + f(B, C)$$




Local scheduler of A: I don't want to do this

```
AxB = matrixMut.remote(A, B)
f_BC = matrixFunction.remote(B, C)
result = matrixSum.remote(AxB, f_BC)
```

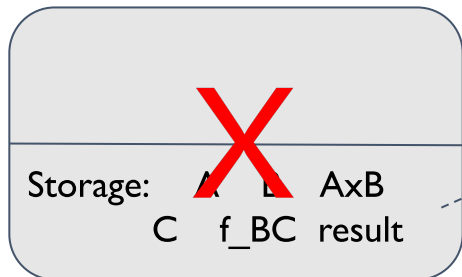


```
result2 = matrixSum(result, A)
```

$$A \times B + f(B, C)$$

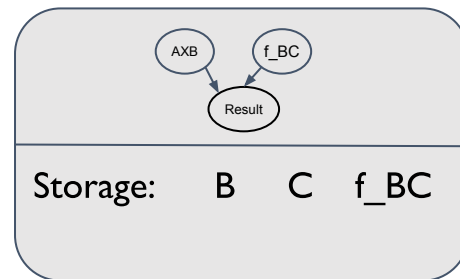


Node A



~~log dependencies~~

Node B



Local scheduler of A: I don't want to do this


```
AxB = matrixMut.remote(A, B)
f_BC = matrixFunction.remote(B, C)
result = matrixSum.remote(AxB, f_BC)
```



```
result2 = matrixSum(result, A)
```

Evaluation

Ray: Locality-aware task placement




Ray: End-to-end scalability


60 nodes

~1,000,000 tasks per sec

The design of GCS &
bottom-up scheduler enables
high horizontal scalability




Evaluation: Allreduce



Ray vs OpenMPI

Scheduler performance is critical




Ray scheduler ablation

Evaluation: RL Application


- The reference system fails to scale to 2048 cores, due the capacity of application driver
- Ray implementation uses an aggregation tree of actors

2x cores ~ average 1.6x speedup




Evolution Strategies


Lineage Stash: Fault tolerance for free?



(a) Deterministic.




(b) Nondeterministic, unlimited forwarding.




(c) Nondeterministic, forward up to 8 nodes.

When uncommitted lineage grows too large, the performance of lineage stash will be greatly reduced

Lineage Stash: Latency during failure



Lineage Stash: Throughput during failure



Summary

Ray: a general-purpose system that

- supports training, serving and simulation efficiently
- unifies stateless (task) and stateful (actor) computations
- has high throughput, low latency and horizontal scalability

Lineage Stash:

asynchronously log the lineage and forward uncommitted lineage to guarantee recovery correctness