

Retiarii: A Deep Learning Exploratory-Training Framework

Authors: Quanlu Zhang, Zhenhua Han, Fan Yang, Yuge Zhang, Zhe Liu,
Mao Yang, Lidong Zhou


Members: Yibo Pi, Jiachen Liu and Qinye Li

Prevalence of Deep Neural Network (DNN)

- Success in perception-based tasks: vision and speech
- Widely used to empower many cloud/edge applications
- Important to design new DNN models

How DNNs Are Commonly Designed?





- Designing new DNNs is an exploratory process



This process is called exploratory-training.

Typical Types of Model Space Explorations


2. Adding a skip connection




1. Replace an operator

3 & 4. Rule-based replacement and insertion,
e.g., inserting *relu* between *conv* and *dense*

Typical Types of Model Space Explorations



5. Generalizing a cell structure to find a better one



Evolving from a base model


Typical Implementation of Model Space Exploration

1. Specifying model space:

a jumbo model with many control flows

2. Exploring model space:

encoding an exploration strategy in the jumbo model

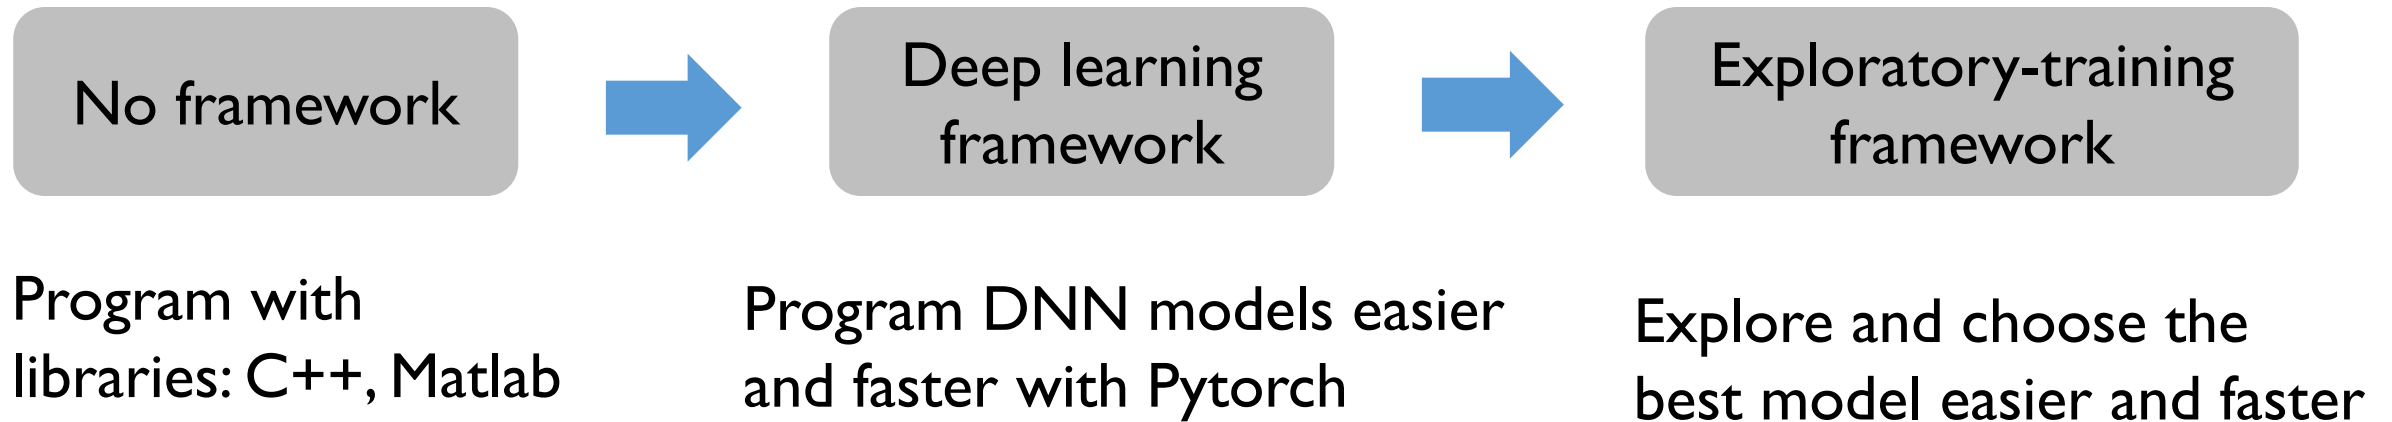


A coupling between model space and exploration strategy

Existing Ways to Achieve Exploratory-Training

- Manually try each new DNN model with deep learning frameworks, e.g., PyTorch
- Neural architecture search (NAS) algorithms or AutoML systems
 - **Lack of modularity**: coupling between model space, exploration strategy, and model training
 - **Lack of reusability**: model space and exploration strategy need to be parameterized
- Missed opportunities to exploit model similarities to speed up the exploration process

A Framework for Exploratory Training?



Retiarii: A Deep Learning Exploratory-Training Framework

- **Key idea**

- Program the model space rather than single models
- model space = base model + mutators (basic operations to modify models)

- **Mutator as the core abstraction**

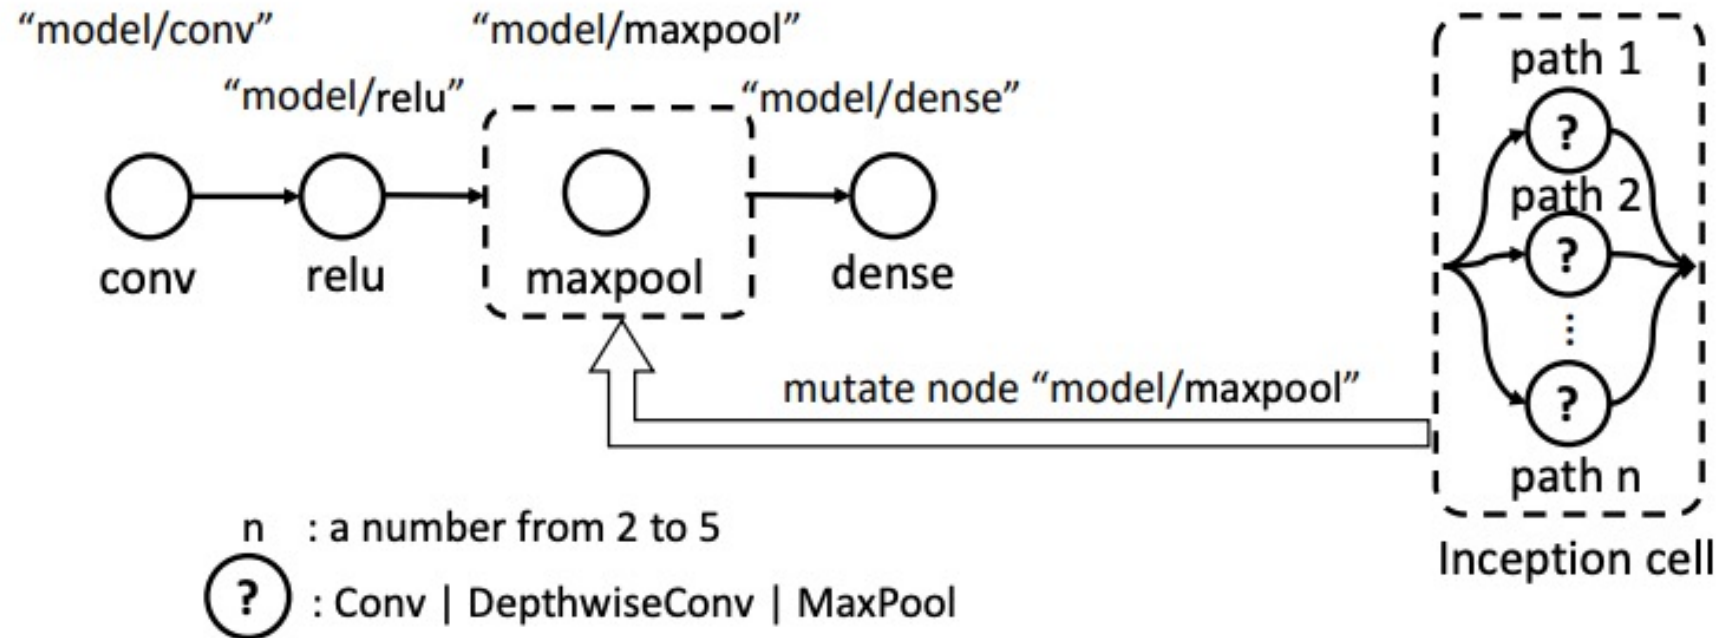
1. Defining arbitrary model space with mutators
2. Decoupling model space from model exploration
 - Model exploration is a set of policies or rules
 - One model exploration strategy can be used for different model spaces
3. Exposing correlations between models
 - Easy to analyze model similarities with mutators

Mutator-based Programming Paradigm

- Mutation primitives

```
1 create_node(name:str, op:Op, params:dict={})
2 delete_node(node:Node)
3 connect(src:NodeOutput, dst:NodeInput)
4 del_connect(src:NodeOutput, dst:NodeInput)
5 update_node(node:Node, op:Op=None, params:dict={},
6             inputs:list=None)
7 choose(candidates:list, n_chosen:int=1,
8        type:str="choice", ctx:dict=None)
```


An Example of Mutating a Base Model



Replace the third node in this 4-node base model with an inception cell

An Example of Mutating a Base Model


- Define the mutator



```
1 # define the graph mutation behavior
2 class InceptionMutator(BaseMutator):
3     def __init__(self, paths_range, candidate_ops):
4         self.paths_range = paths_range # [2, 3, 4, 5]
5         self.ops = candidate_ops # {conv, dconv, ...}
6     def mutate(self, targets):
```

An Example of Mutating a Base Model

- Check the target node chain (relu, maxpool, dense)




```
1 # define the graph mutation behavior
2 class InceptionMutator(BaseMutator):
3     def __init__(self, paths_range, candidate_ops):
4         self.paths_range = paths_range # [2, 3, 4, 5]
5         self.ops = candidate_ops # {conv, dconv, ...}
6     def mutate(self, targets):
7         if not three_node_chain(targets):
8             return err
```

An Example of Mutating a Base Model

- Choose # of paths

“model/conv”



“model/relu”



“model/dense”




Two paths

```
1 # define the graph mutation behavior
2 class InceptionMutator(BaseMutator):
3     def __init__(self, paths_range, candidate_ops):
4         self.paths_range = paths_range # [2, 3, 4, 5]
5         self.ops = candidate_ops # {conv, dconv, ...}
6     def mutate(self, targets):
7         if not three_node_chain(targets):
8             return err
9         n = choose(candidates=self.paths_range)
10        delete_node(targets[1])
```

An Example of Mutating a Base Model

- Creating paths

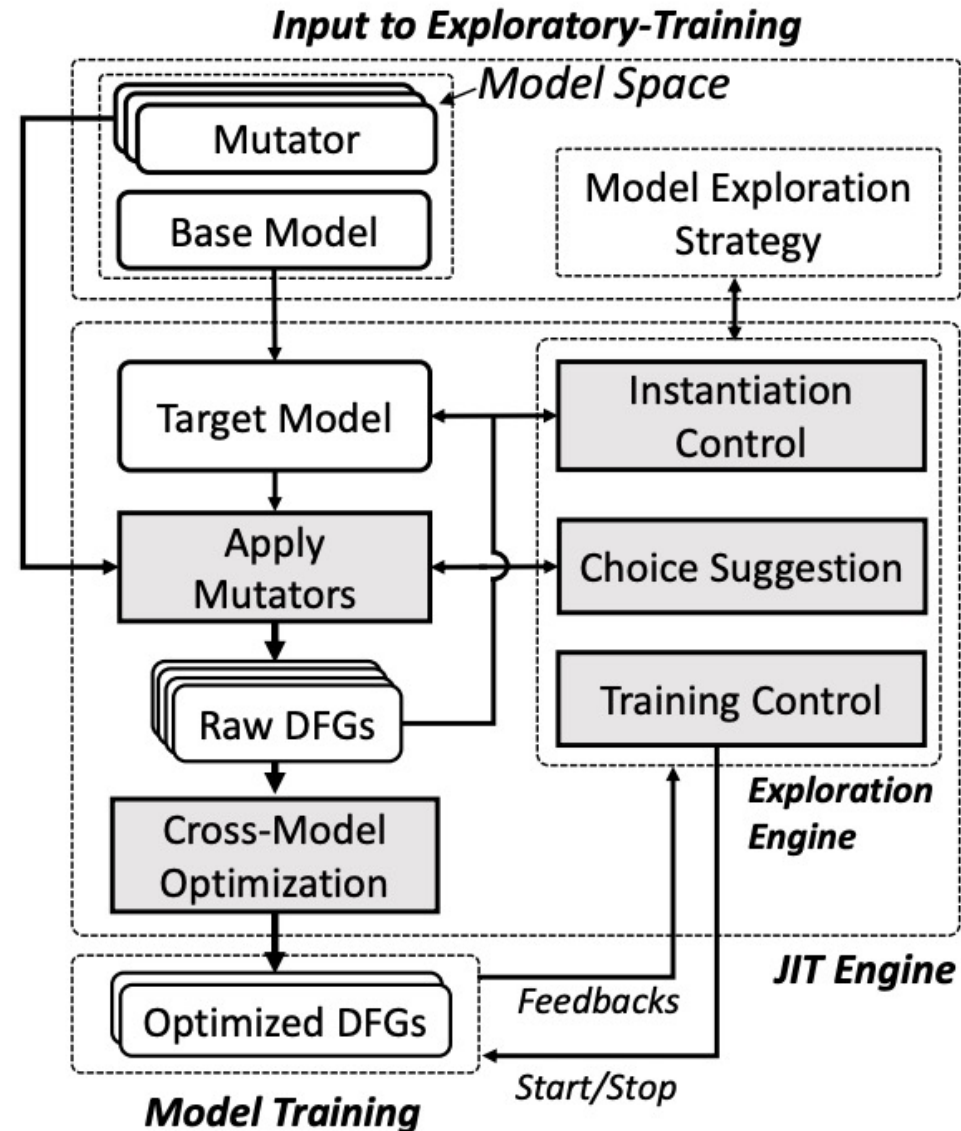


```
1 # define the graph mutation behavior
2 class InceptionMutator(BaseMutator):
3     def __init__(self, paths_range, candidate_ops):
4         self.paths_range = paths_range # [2, 3, 4, 5]
5         self.ops = candidate_ops # {conv, dconv, ...}
6     def mutate(self, targets):
7         if not three_node_chain(targets):
8             return err
9         n = choose(candidates=self.paths_range)
10        delete_node(targets[1])
11        for i in range(n): # create n paths
12            op = choose(candidates=self.ops)
13            nd = create_node(name='way_'+str(i), op=op)
14            connect(src=targets[0].output, dst=nd.input)
15            connect(src=nd.output, dst=targets[2].input)
```


System Architecture

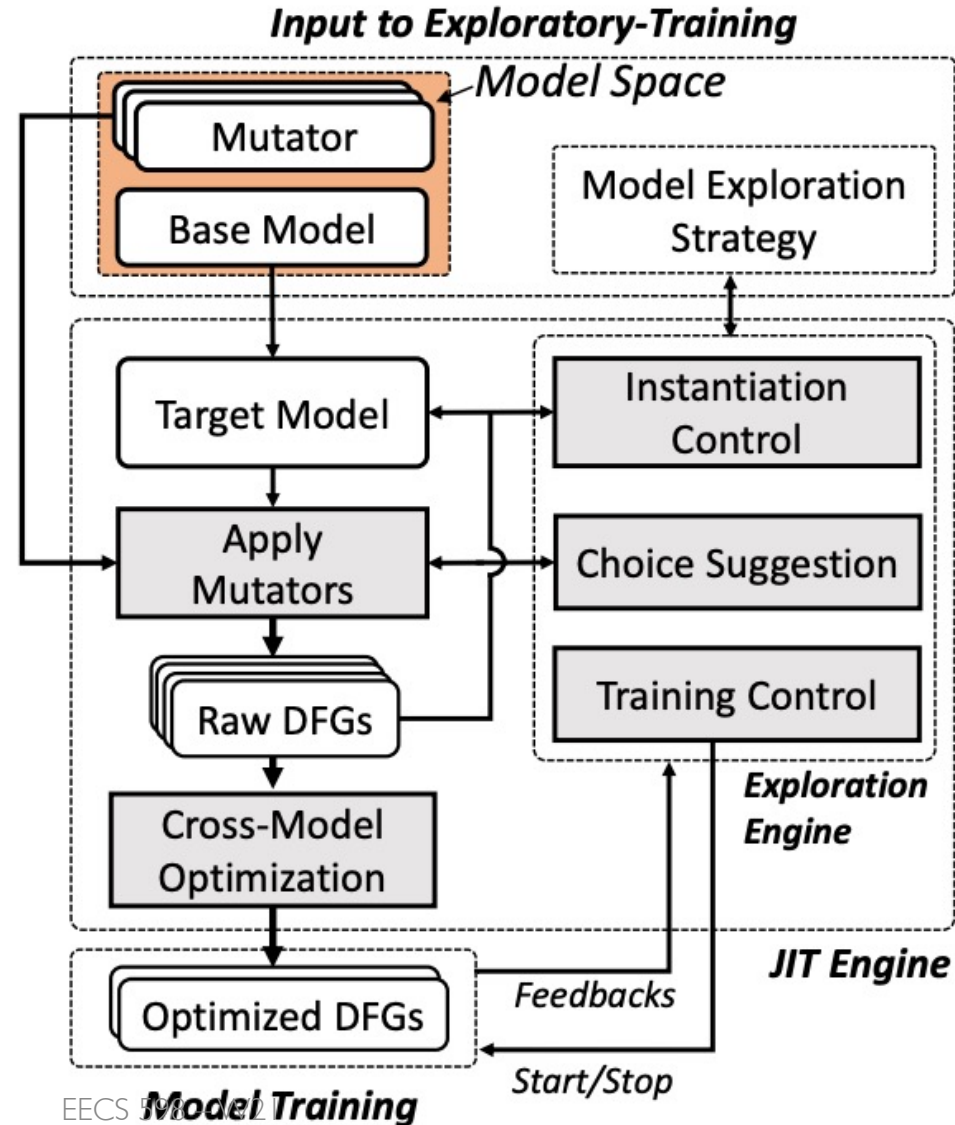
- **Inputs**

- One or more base models
- A set of mutators
- A policy describing the exploration strategy



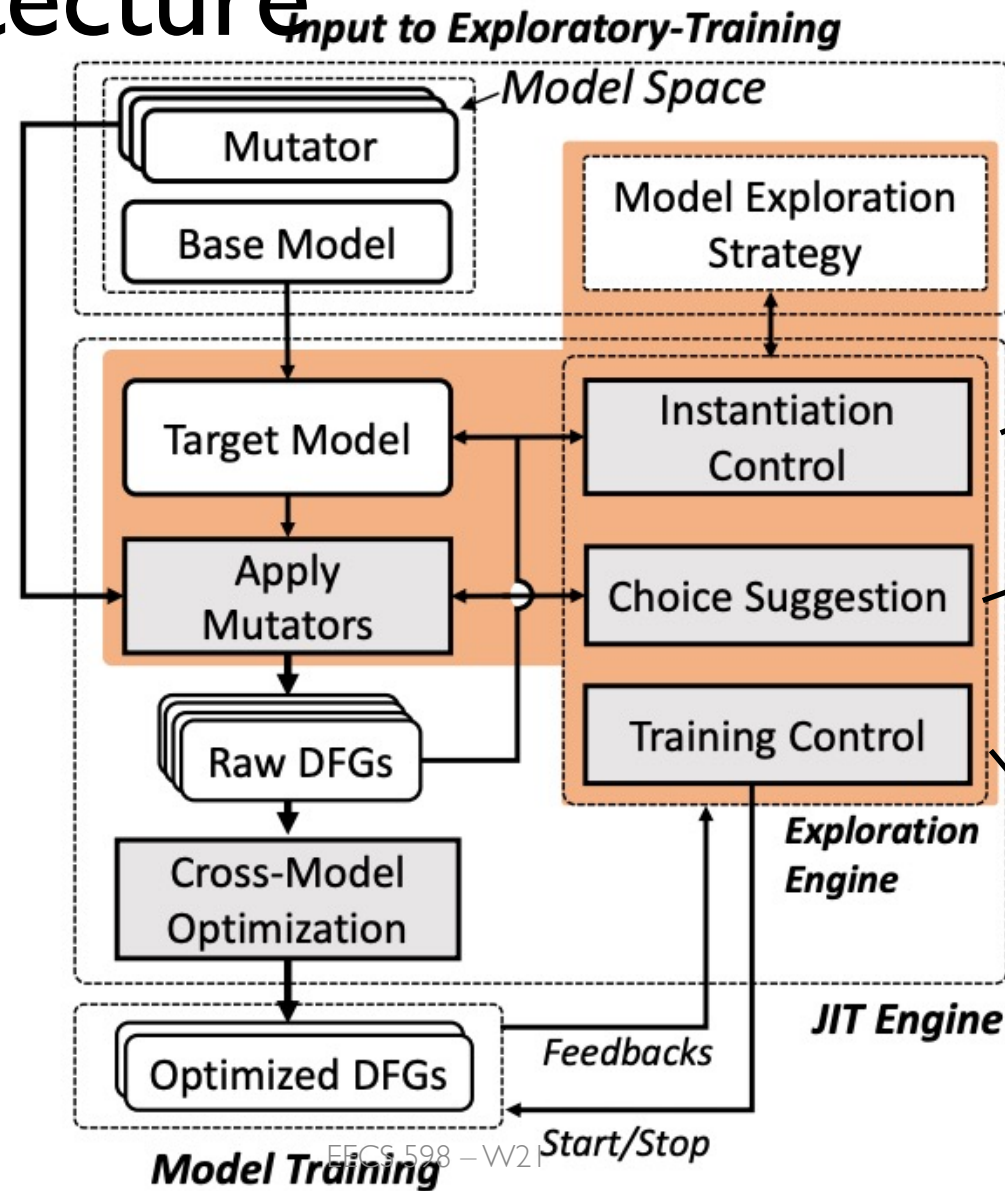
System Architecture

- Model space
 - Mutator
 - Base models



System Architecture

- **Exploration engine**
 - Instantiation control
 - Choice suggestion
 - Training control



Which target model
& mutators to use

Strategy for different
combinations of
mutators

Monitor training,
collect feedbacks, and
control priorities and
resource allocation

- **Model training**
 - Optimized data-flow graphs



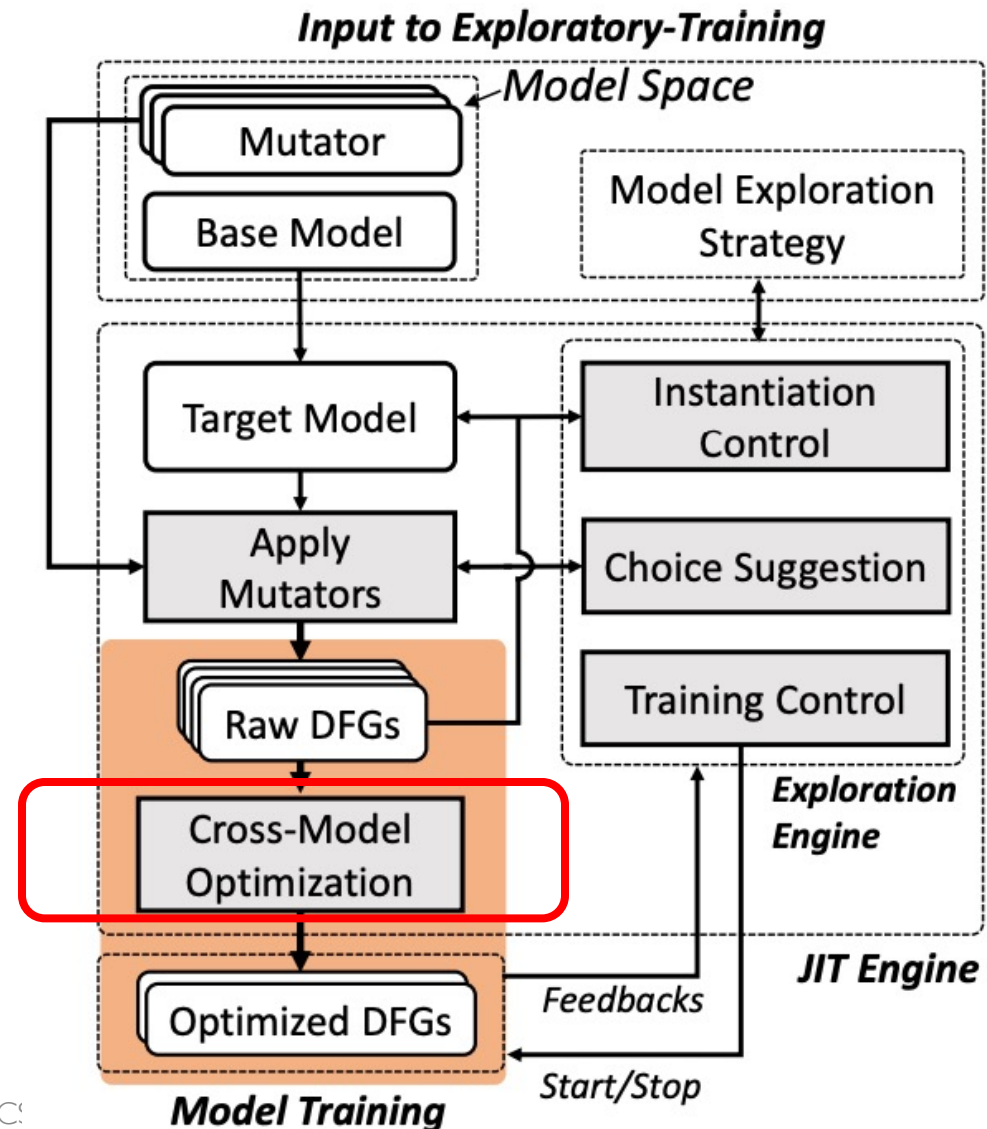
Expressiveness and Reusability

- Retiarii currently supports 27 neural architecture search (NAS) solutions

NAS Solution	Model Space	Exploration Strategy	Required Mutator Class			
			<i>Input Mutator</i>	<i>Operator Mutator</i>	<i>Inserting Mutator</i>	<i>Customized Mutator</i>
MnasNet [59]	MobileNetV2-based space	Reinforcement Learning		✓	✓	
NASNet [70]	NASNet cell	Reinforcement Learning	✓	✓		
ENAS-CNN [50]	NASNet cell variant	Reinforcement Learning	✓	✓		
AmoebaNet [51]	NASNet cell	Evolutionary	✓	✓		
Single-Path One Shot (SPOS) [27]	ShuffleNetV2-based space	Evolutionary		✓		
Weight Agnostic Networks [23]	Evolving space w/ adding/altering nodes adding connections	Evolutionary		✓		✓
Path-level NAS [13]	Evolving space w/ replication and split	Reinforcement Learning				✓
TextNAS [62]	TextNAS space	Reinforcement Learning	✓	✓		
...

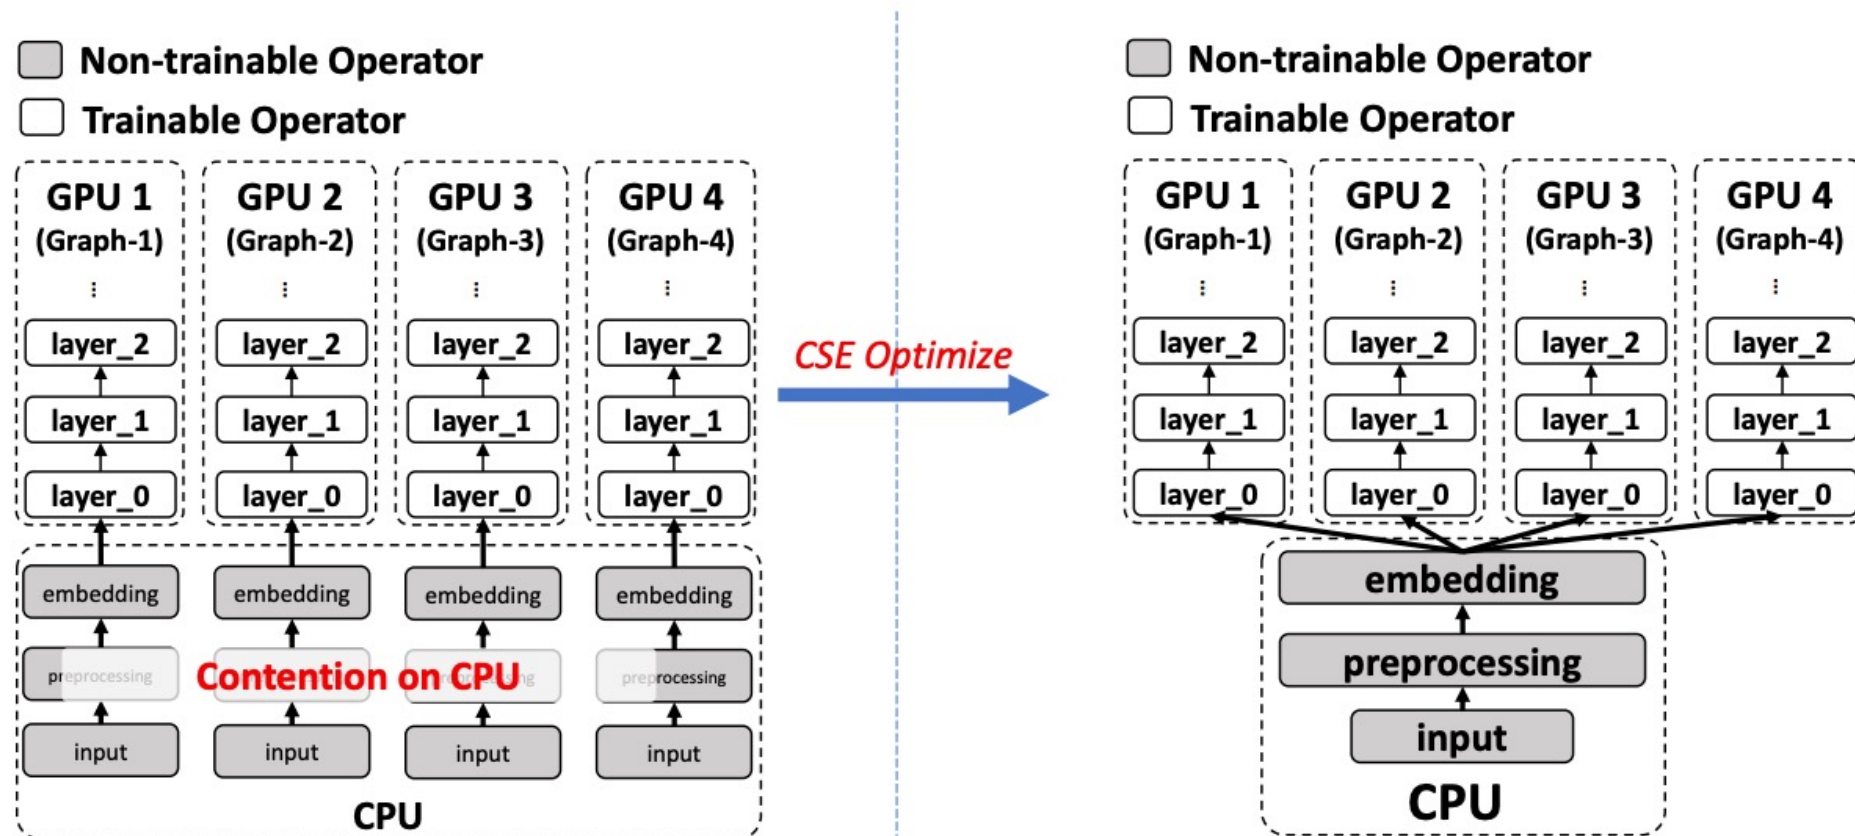
Cross-Model Optimization

- Optimization opportunities
 - Same training data
 - Same data preprocessing
 - Common layers
 - Weight sharing among models




Common Sub-expression Elimination (CSE)

- De-duplicating CPU-based common prefix operations



Common Sub-expression Elimination (CSE)

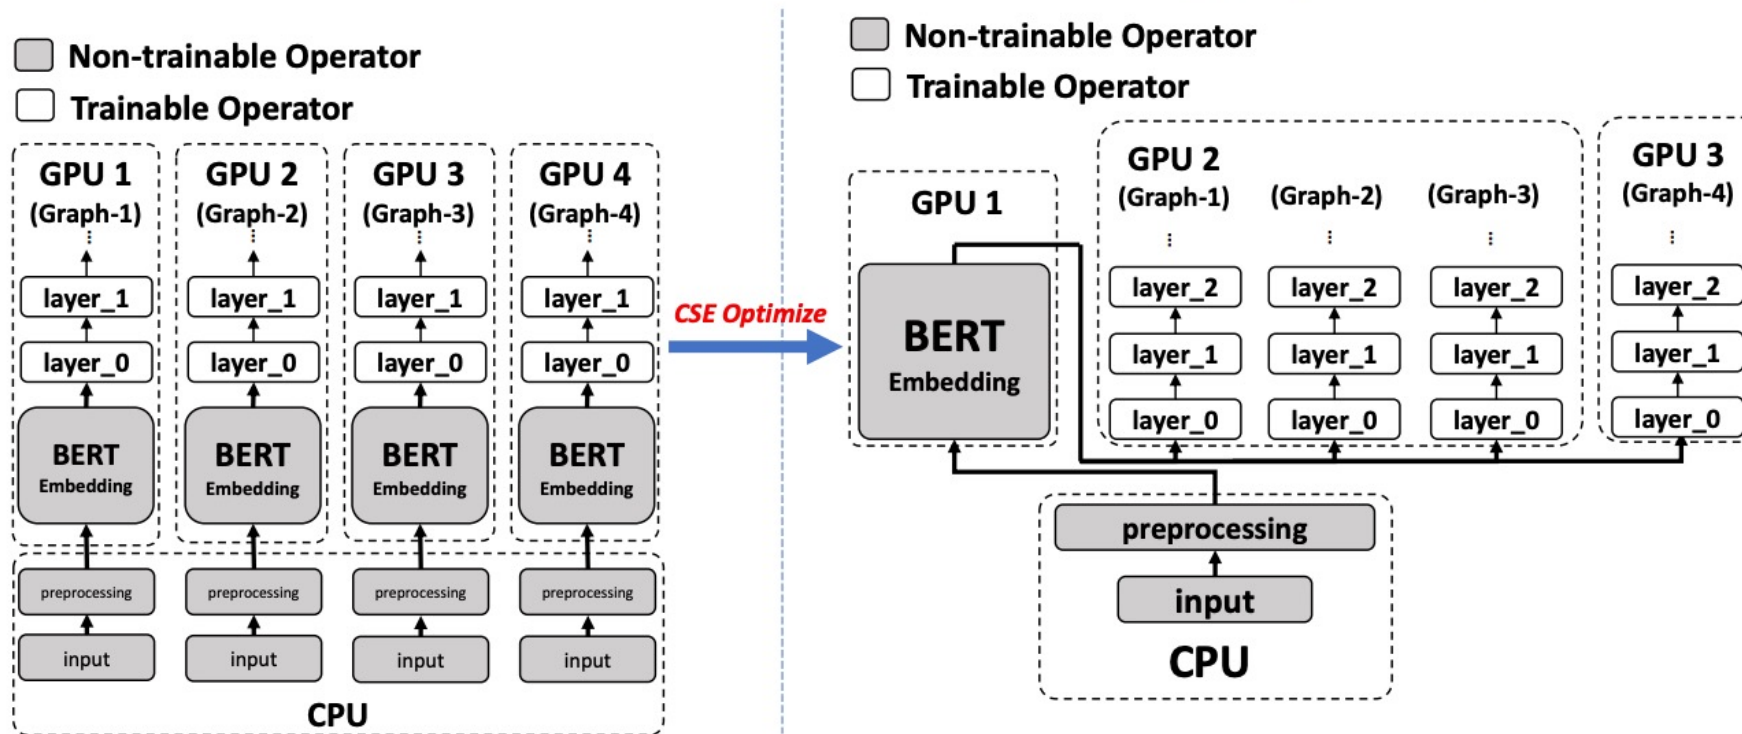
Experiment: training MnasNet0.5 on ImageNet with 4 V100 GPUs and 20 CPU cores



CSE of CPU-based operation

Common Sub-expression Elimination (CSE)

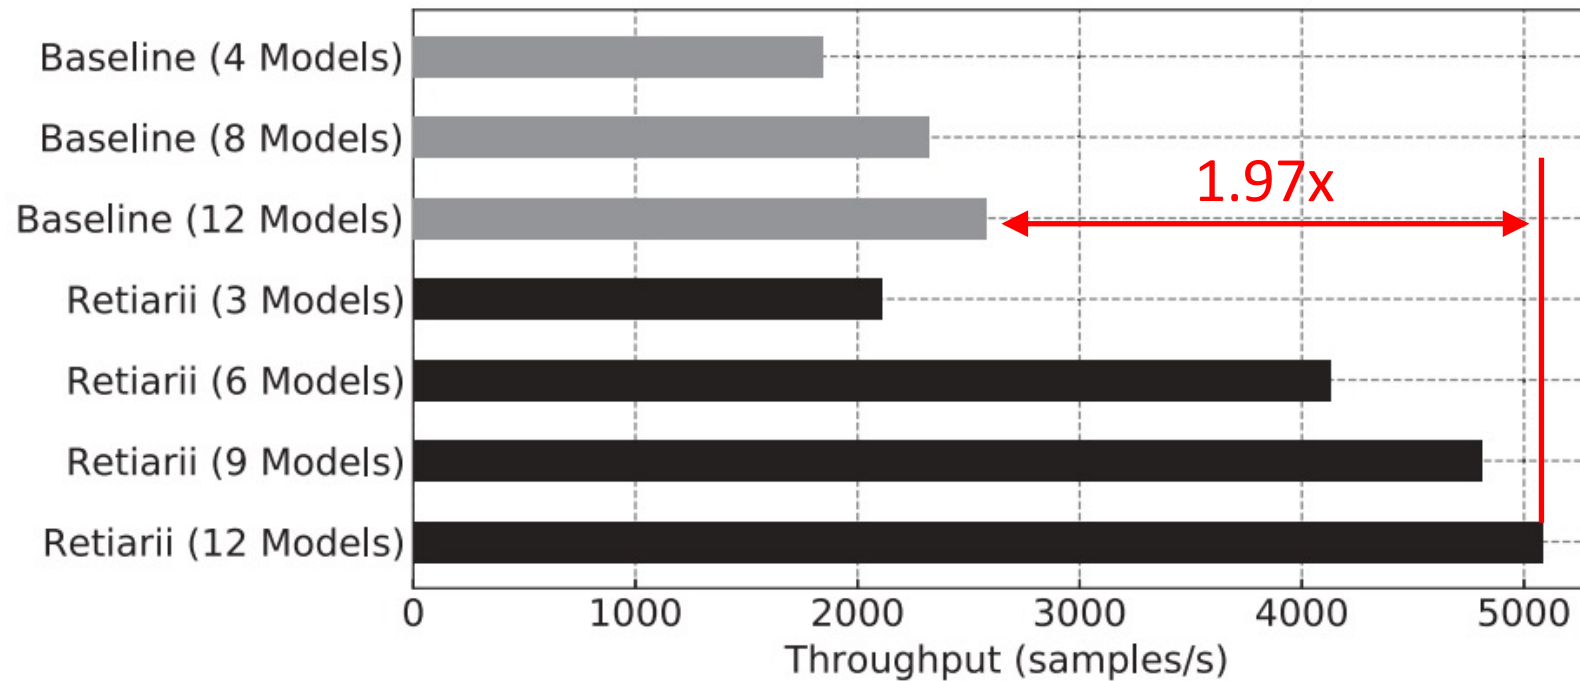
- CSE + device placement for GPU-based embedding



Dedicating one GPU for BERT-embedding improves pipeline and reduce memory consumption

Common Sub-expression Elimination (CSE)

Experiment: training TextNAS, one of the state-of-the-art NLP models



Speeding up Neural Architecture Search (NAS)

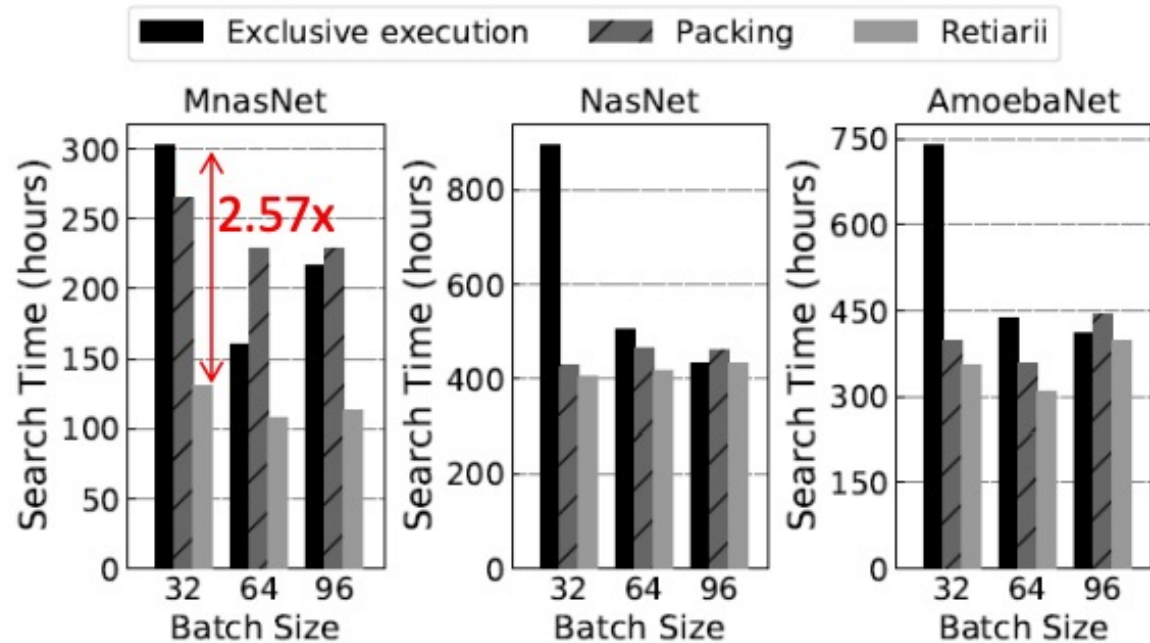
- Three popular NAS solutions

NAS Solution	Search Space	Exploration Strategy
MnasNet	Factorized Hierarchical Search Space	Reinforcement Learning
NASNet	Normal Cell + Reduction Cell	Reinforcement Learning
AmoebaNet	Normal Cell + Reduction Cell	Evolutionary Algorithm


- **Baselines**
 - Exclusive execution: training one model per GPU at a time
 - Packing: training multiple models per GPU using NVIDIA CUDA MPS

Speeding up Neural Architecture Search (NAS)

- Running on 4 NVIDIA Tesla V100 GPUs
- Training 1000 models for 1 epoch on ImageNet



(a) NVIDIA Data Loading Library (DALI)



(b) PyTorch DataLoader

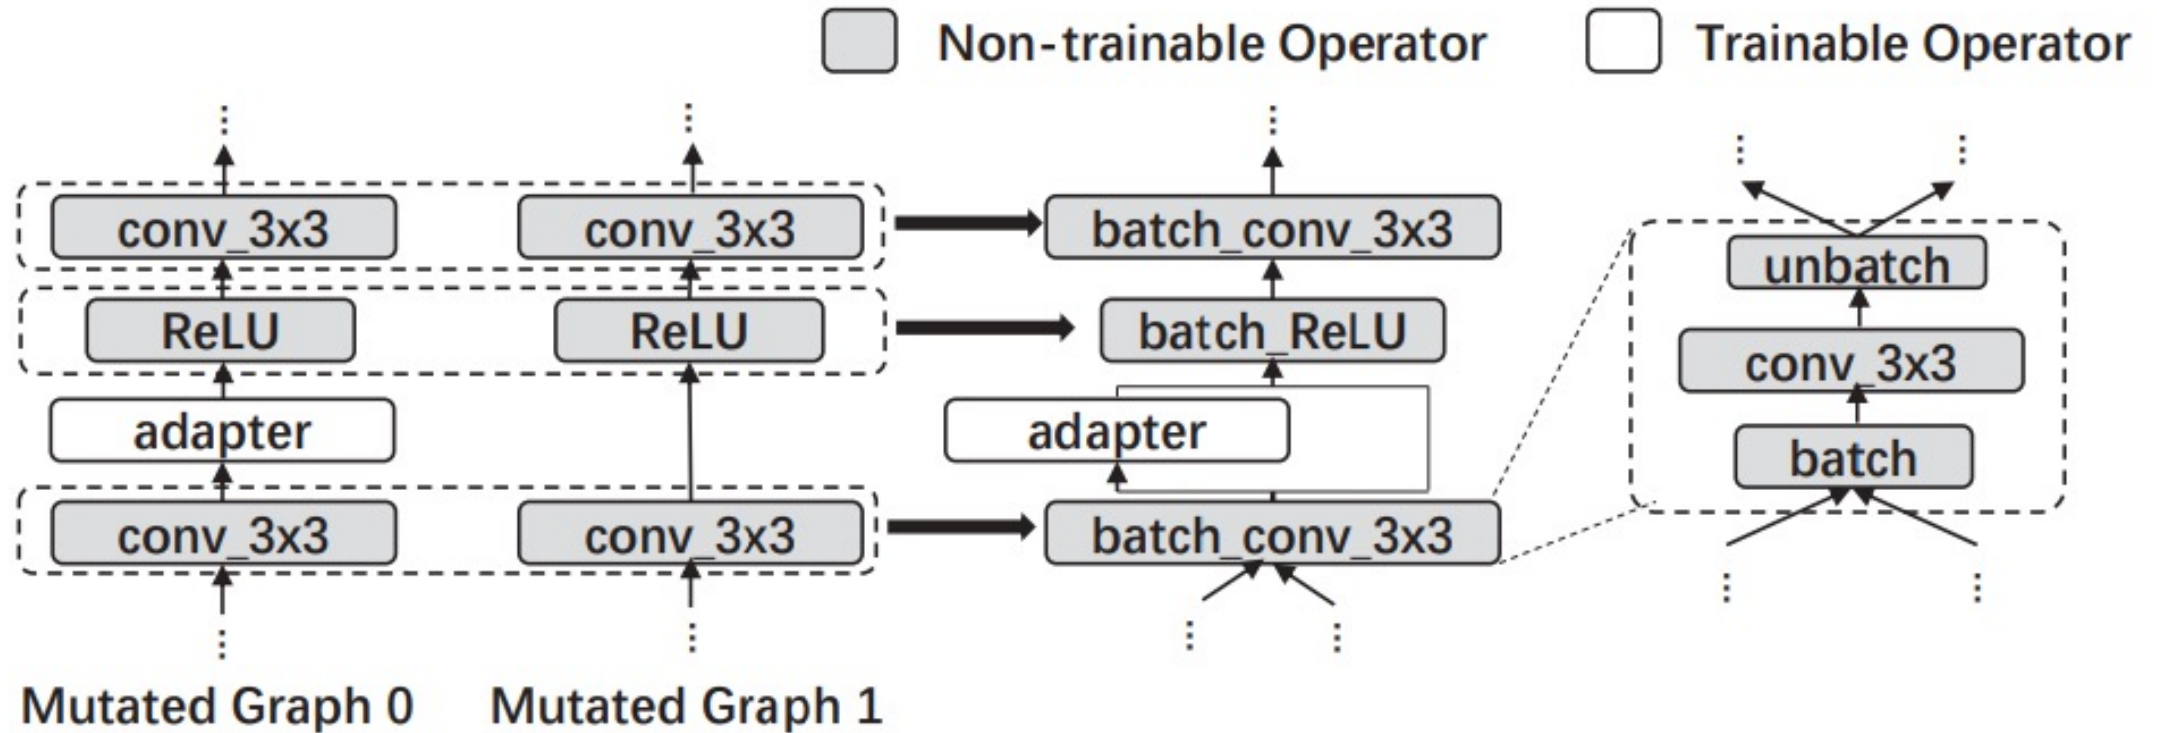
Conclusion

- Retiarri is a new DNN framework for exploratory training
- Retiarri provides new interfaces for model developers to explore new models efficiently
- Retiarri uses the Mutator abstraction to achieve
 - Strong expressiveness in model space
 - Reusability of exploration strategies
 - Cross-model optimization

Discussion


- **Limitations of Retiarii mentioned by the authors**
 1. Limited support to dynamic graphs
 2. Limited support to operator batching
 3. Possible shape mismatch between adjacent layers' input/output tensors
- **Other limitations:**
 1. A whitelist is used to identify operators requiring dedicated GPUs
 2. Retiarii greedily packs as many models as possible into one GPU
 - What if single models are too large to fit in to one GPU?
 3. What if there are significant mutations to the base model?
 - Is it easy for developers to manage these mutations?
 - How will cross-model optimization perform when models are very different?

Operator Batching




Speed up Weight-Shared Training


- What is weight sharing?



Trial #1




Trial #2




Speed up Weight-Shared Training

- Building a super-graph to encode the search space



Speed up Weight-Shared Training


- Building a super-graph to encode the search space



Limited space size and hard to scale to a large GPU cluster

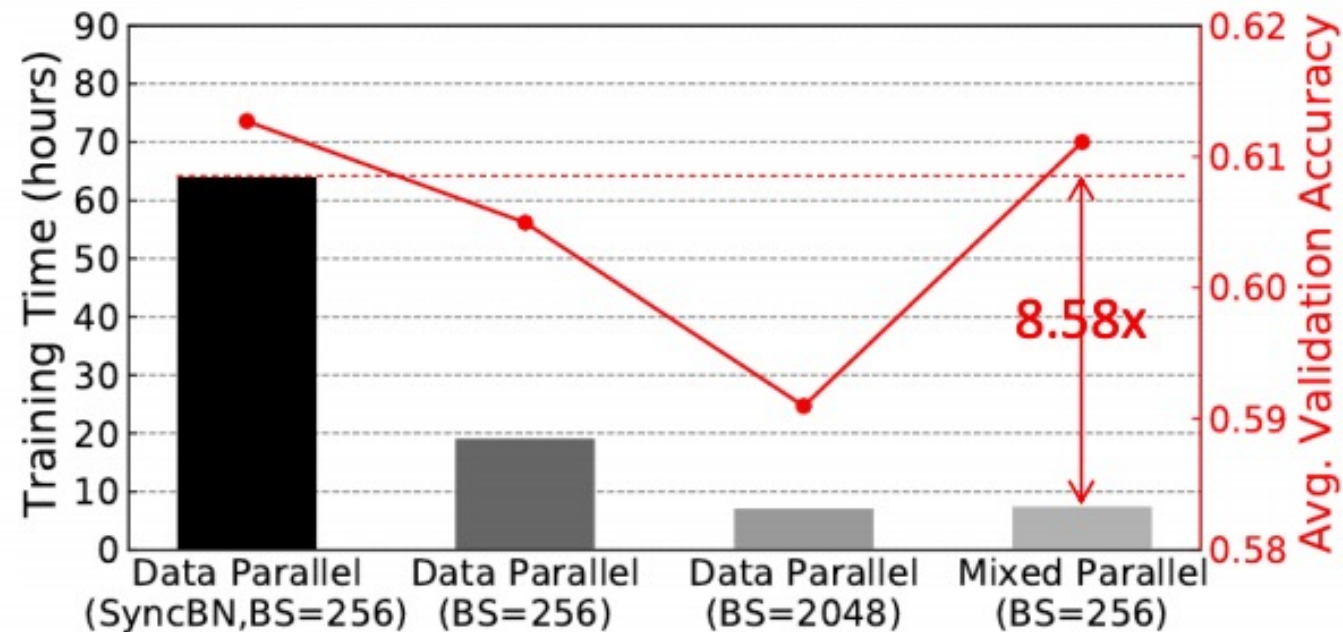
Speed up Weight-Shared Training

- **Mixed Parallelism for weight sharing**
 - Model parallelism partitions the super-graph to multiple GPUs
 - Data parallelism feeds each partition with a different batch of data



Speed up Weight-Shared Training

- Experiment with a popular weight-shared NAS, SPOS [*]



[*] Guo Z, et al. “Single path one-shot neural architecture search with uniform sampling”. arXiv preprint. 2019 Mar 31.

Fluid: Resource-aware Hyperparameter Tuning Engine

Peifeng Yu[†], Jiachen Liu[†], Mosharaf Chowdhury

[†] Equal
contribution



SymbioticLab



Outline

1. Background and Motivation

2. Abstraction and Algorithms

3. Evaluation



Hyperparameter Tuning Today

- Hyperparameters
 - # of layers/# of neurons
 - Dropout rate
 - # of channels
 - Learning rate
 - Optimizer parameters
 - Etc.
- Non-differentiable & high dimensional search space

Hyperparameter Tuning Today


- Evaluation of hyperparameters is time/resource consuming
 - train a model to know if it works
- Many algorithms & techniques
 - Random/Grid
 - Model-based config generation (BO, PTE ...)
 - Bandit-based / early-stopping (SHA, HB, BOHB...)
 - Many others (ASHA...)

Example: SHA




1. Resource underutilized
2. Straggler problem

Example: ASHA / BOHB



1. Maximize resource utilization by improving parallelism

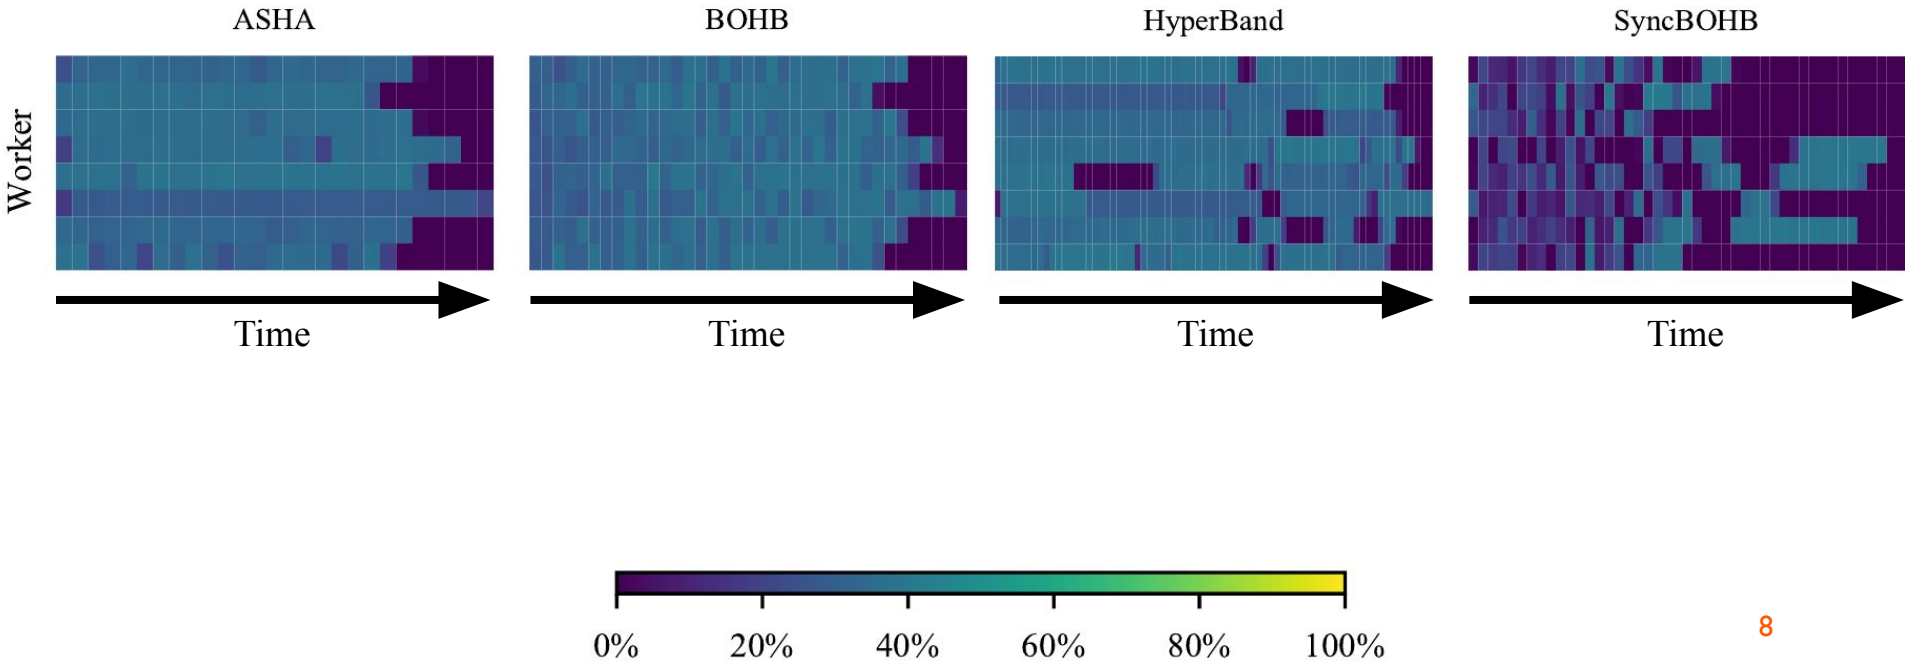
Example: BOHB



1. Higher time to accuracy with the increase of #worker
2. Unknown resource to accuracy performance (goodput)
3. Simply improving parallelism may waste resources

Trials Execution

Resource utilization overtime



Resource Parallelism

Inter-GPU parallelism: Distributed training

- Fully utilize the idle resources

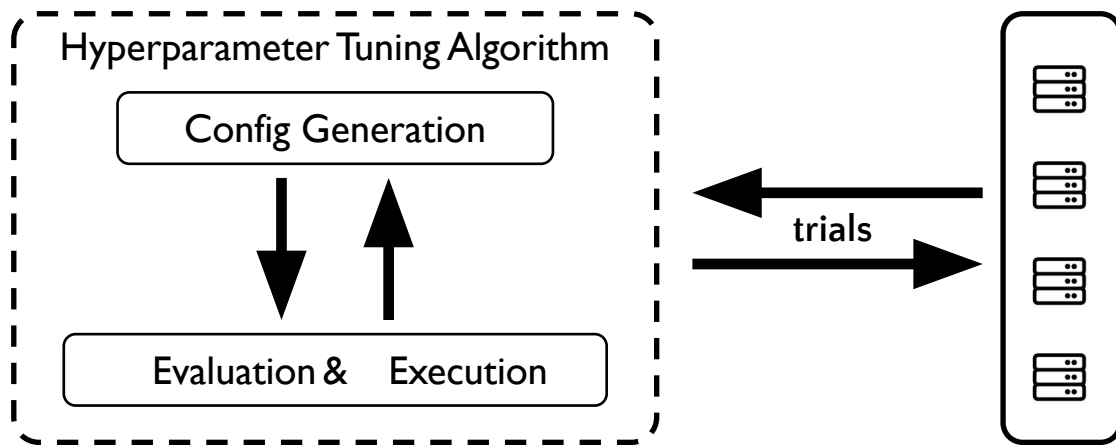
Intra-GPU parallelism: Nvidia MPS

- Execute more trials with under-utilized resources

Goal:

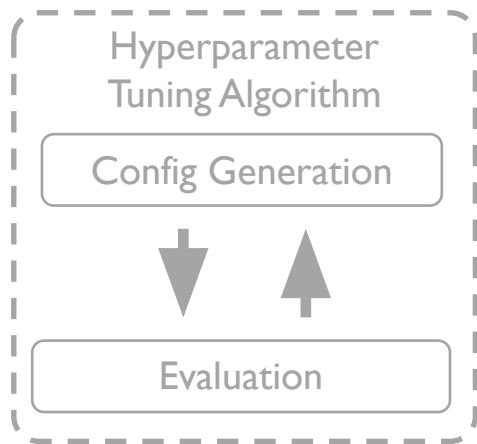
1. Improve resource utilization
2. Minimize the makespan

Hyperparameter Execution Engine: Fluid




- **Direct** interaction with the cluster to execute **trials**
- Trials gets executed in FIFO order


Hyperparameter Execution Engine: Fluid



Execution



Hyperparameter Execution Engine: **Fluid**



Challenge:

- Wide variety of tuning algorithms
 - Random/Iterative/Sequential
 - ✓ **TrialGroup**
- Heterogeneity & dynamicity
- ✓ Integrated algorithm for leveraging multiple source of parallelism

Outline

1. Background and Motivation

2. Abstraction and Algorithms


3. Evaluation

The Interface: TrialGroup

- Definition


A group of training trials with a training budget associated to each trial.

- Example

- Given 5 trials to evaluate:  x5

- Grid/random search:

- Sequential model-based algorithms:



The Interface: TrialGroup

- Definition

A group of training trials with a training budget associated to each trial.

- Generalization

- All kinds of hyperparameter tuning algorithms could be expressed by **a sequence of TrialGroup** and executed by Fluid.

Problem Definition: Strip Packing

- Input: TrialGroup $A = \{a_1, a_2, \dots, a_k\}$, resources $M = \{m_1, m_2, \dots, m_n\}$
- Output: resource allocation $W = \{w_1, w_2, \dots, w_n\}$
- Goal: minimize the length L of strips

$w = 1$



$l = 30$


$w = 1$



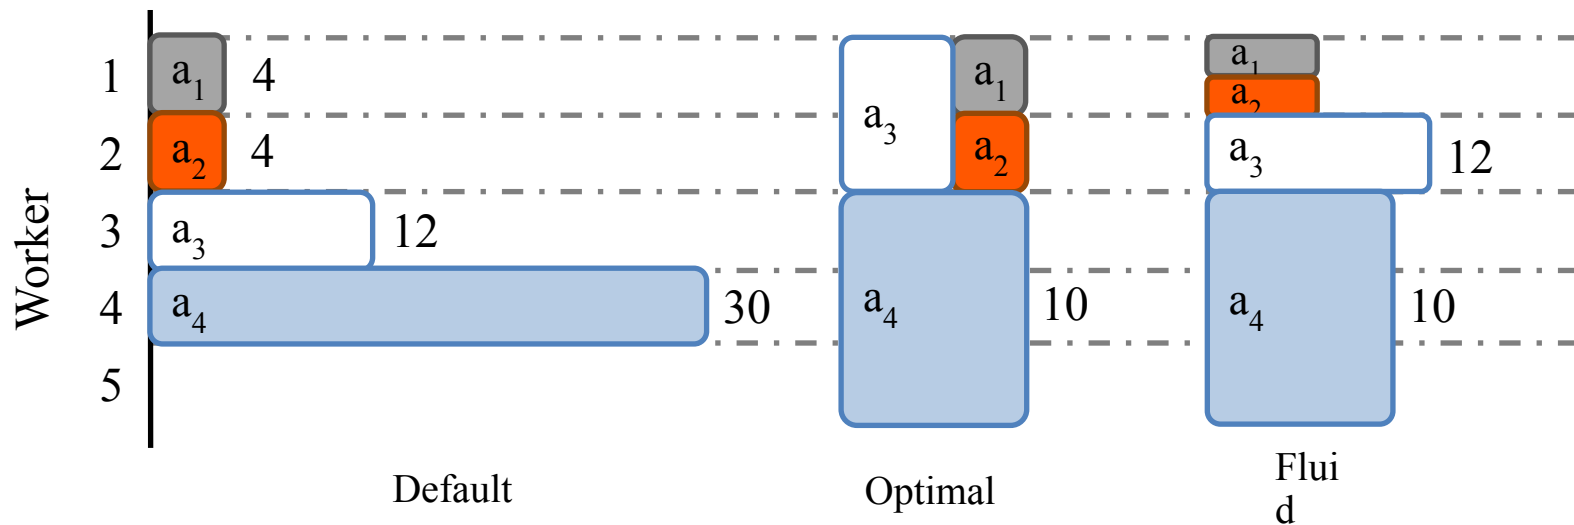
$l = 1$

Problem Definition: Strip Packing

- Input: TrialGroup $A = \{a_1, a_2, \dots, a_k\}$, resources $M = \{m_1, m_2, \dots, m_n\}$
- Output: resource allocation $W = \{w_1, w_2, \dots, w_n\}$
- Goal: minimize the length L of strips



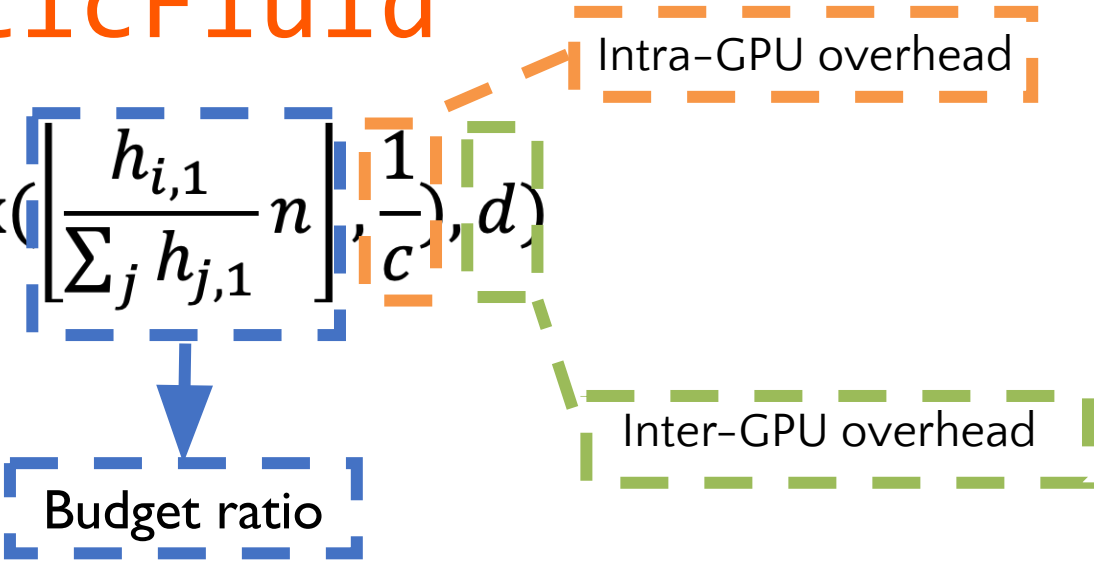
Toy Example



Different solutions to execute 4 trials (1 TrialGroup) scheduled on 5 workers


Fully **utilize the resources** and **mitigate the straggler**

Algorithm: StaticFluid

$$w_i = \min(\max(\left\lfloor \frac{h_{i,1}}{\sum_j h_{j,1}} n \right\rfloor, \frac{1}{c}), d)$$


- h : trial training budget
- n : available resources
- c : maximum intra-GPU parallelism (# of packing trials)
- d : maximum inter-GPU parallelism (# of distributed workers)

Algorithm: DynamicFluid



Fluid is event-driven:

1. Trials added / removed
2. Resource added / changed

Algorithm: DynamicFluid

$$w_i = \min(\max(\left\lfloor \frac{h_{i,1}}{\sum_j h_{j,1}} n \right\rfloor, \frac{1}{c}), d)$$

if $w'_i > w_i$ and $h_{i,w'_i} + \epsilon < h_{i,w_i}$

Update a_i with w_i resources \triangleright Scale up

else if $w'_i < w_i$ and $w'_i(h_{i,w'_i} + \epsilon) < w_i h_{i,w_i}$

Update a_i with w_i resources \triangleright Scale down

ϵ : scale up / down overhead

Outline

1. Background and Motivation

2. Abstraction and Algorithms

3. Evaluation


Evaluation Setup

- Implementation: an alternative Ray^[1] executor
- Workloads


Task	Base Model	# of Params.	Target
CIFAR-10	AlexNet	7	Acc. \geq 90%
WLM	RNN	10	PPL \leq 140
DCGAN	CNN	2	Inception \geq 5.2

Evaluation Results

- Average resource utilization: 10%-100% improvement
- Average job completion time: 10%-70% improvement



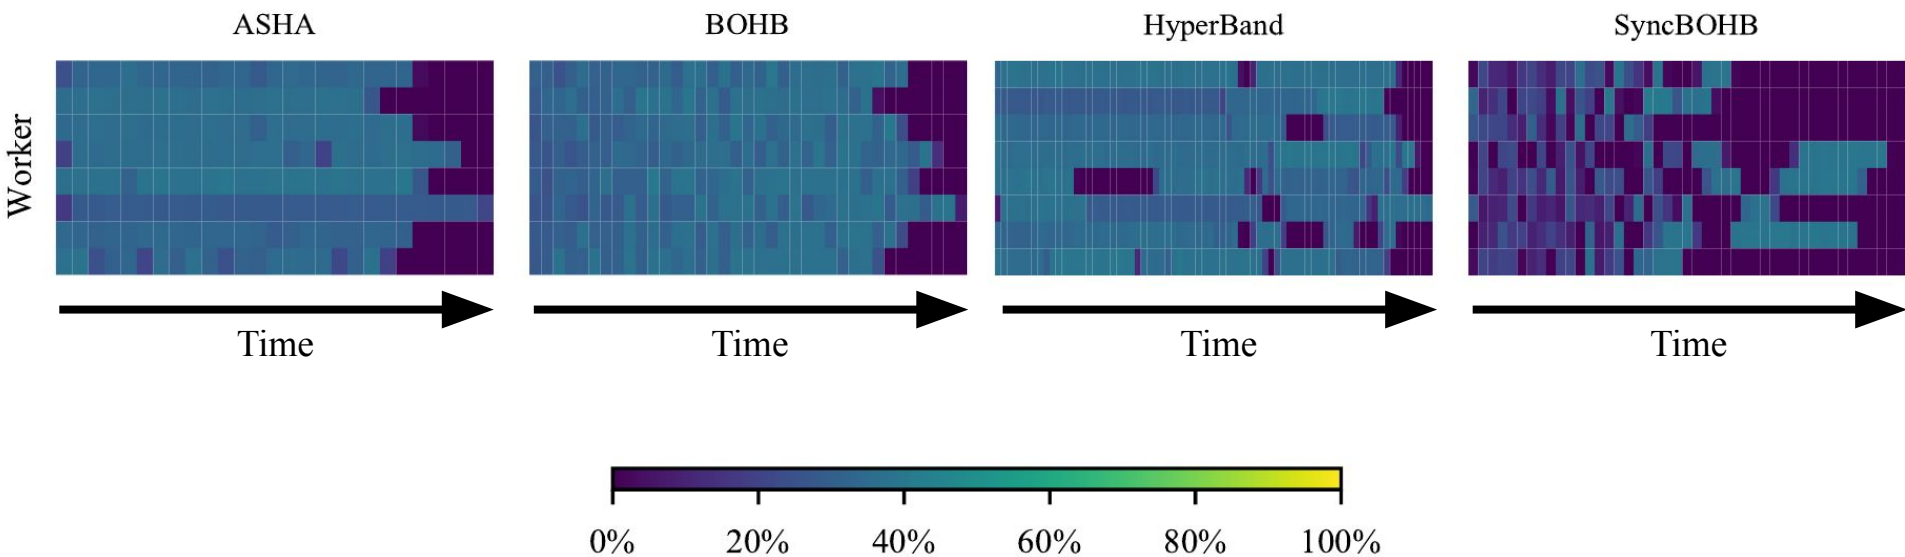
(a) CIFAR-10



(b) CIFAR-10

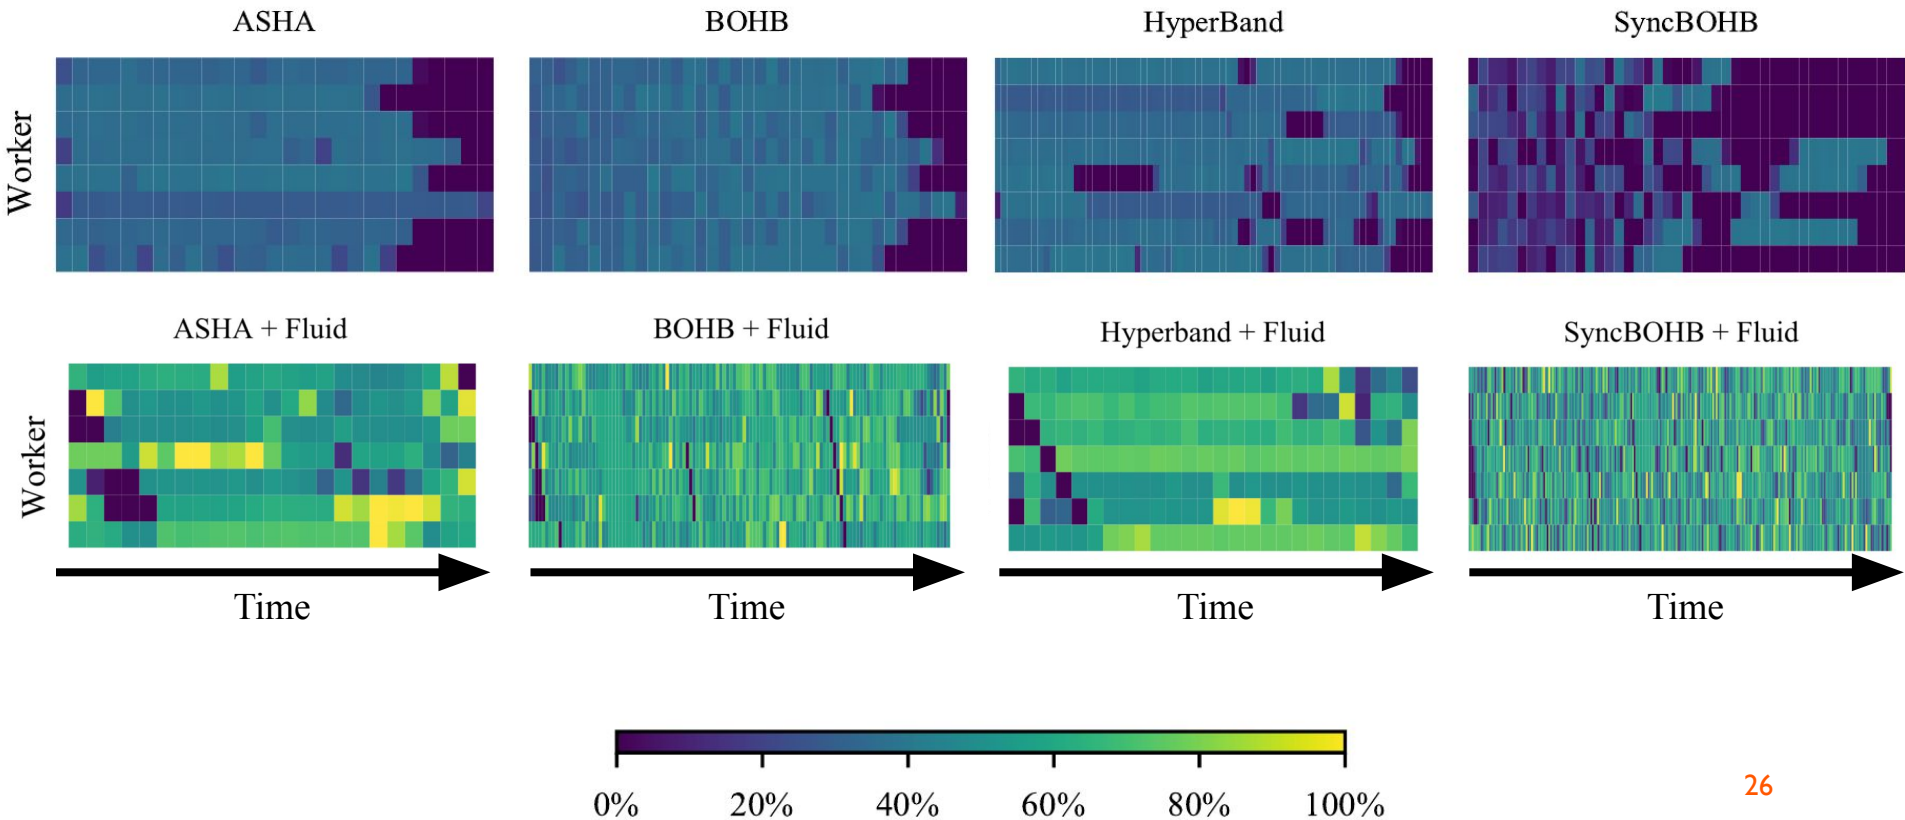
Evaluation Results: Visualization

Resource utilization over time



Evaluation Results: Visualization

Resource utilization over time



Conclusion

- Fluid
 - Hyperparameter tuning execution engine
 - Can be combined with most tuning algorithms
 - Improve utilization and end-to-end tuning time
- Open source
 - <https://github.com/SymbioticLab/fluid>
- Q&A

